

Array and Hashing

1. Two Sum

Pattern: Arrays & Hashing

Problem Statement

Given an array of integers `nums` and an integer `target`, return indices of the two numbers such that they add up to `target`.

You may assume that each input would have exactly one solution, and you may not use the same element twice.

You can return the answer in any order.

Sample Input & Output

```
Input: nums = [2,7,11,15], target = 9
```

```
Output: [0,1]
```

```
Explanation: nums[0] + nums[1] == 9, so we return [0, 1].
```

```
Input: nums = [3,3], target = 6
```

```
Output: [0,1]
```

```
Explanation: Two identical elements at different indices are valid.
```

Input: nums = [1,2,3], target = 7
Output: [] (or raises; but problem guarantees one solution)

LeetCode Editorial Solution + Inline Tests

```
from typing import List

class Solution:
    def twoSum(self, nums: List[int], target: int) -> List[int]:
        # STEP 1: Initialize hash map to store value -> index
        # - Why? To check in O(1) if complement (target - num) exists
        seen = {}

        # STEP 2: Iterate through array with index
        # - Why index? We need to return positions, not values
        for i, num in enumerate(nums):
            complement = target - num

            # STEP 3: Check if complement already seen
            # - If yes, we found our pair: current index + stored index
            if complement in seen:
                return [seen[complement], i]
                # return [complement, num] -> returns the actual number

            # STEP 4: Store current number and index for future lookup
            # - Why here? To avoid using same element twice
            seen[num] = i

        # STEP 5: Return empty if no solution (per constraints, won't happen)
        # - Included for safety / clarity
        return []

# ----- INLINE TESTS -----
if __name__ == "__main__":
    sol = Solution()

    # Test 1: Normal case
    result1 = sol.twoSum([2, 7, 11, 15], 9)
```

```

print(f"Test 1: {result1} → Expected: [0, 1]")
assert result1 == [0, 1], "Test 1 Failed"

# Test 2: Edge case - duplicate values
result2 = sol.twoSum([3, 3], 6)
print(f"Test 2: {result2} → Expected: [0, 1]")
assert result2 == [0, 1], "Test 2 Failed"

# Test 3: Tricky - negative numbers
result3 = sol.twoSum([-1, -2, -3, -4, -5], -8)
print(f"Test 3: {result3} → Expected: [2, 4]")
assert result3 == [2, 4], "Test 3 Failed"

print(" All inline tests passed!")

```

How to use: Copy-paste this block into .py or Quarto cell → run directly → instant feedback.

Example Walkthrough

Let's walk through `nums = [2, 7, 11, 15]`, `target = 9`.

Initial state:

`seen = {}` — empty hash map.

We'll iterate with index `i` and value `num`.

Step 1 — `i=0`, `num=2`:

- `complement = 9 - 2 = 7`

- Is 7 in `seen`? No → skip return

- Store `seen[2] = 0` → `seen = {2: 0}`

→ *Why store?* So if later we see 7, we know 2 was at index 0.

Step 2 — i=1, num=7:

- complement = 9 - 7 = 2
- Is 2 in **seen**? Yes → at index 0
- Return [**seen**[2], 1] → [0, 1]

→ *Why not store 7 first?*

Because we check *before* storing — this ensures we never use same index twice.

→ *Pattern insight:*

We trade space (hash map) for time — instead of nested loops $O(n^2)$, we do one pass $O(n)$. Hashing lets us “remember” what we’ve seen and instantly find complements.

Complexity Analysis

- **Time Complexity:** $O(n)$

We traverse the list once. Each hash map lookup and insertion is $O(1)$ average case.

- **Space Complexity:** $O(n)$

In worst case, we store $n-1$ elements in the hash map before finding the solution.

2. Contains Duplicate

Pattern: Arrays & Hashing

Problem Statement

Given an integer array **nums**, return **true** if any value appears at least twice in the array, and return **false** if every element is distinct.

Sample Input & Output

```
Input: nums = [1,2,3,1]
Output: true
Explanation: The number 1 appears twice.
```

```
Input: nums = [1,2,3,4]
Output: false
Explanation: All elements are unique.
```

```
Input: nums = [1]
Output: false
Explanation: Single element cannot have duplicates.
```

LeetCode Editorial Solution + Inline Tests

```
from typing import List

class Solution:
    def containsDuplicate(self, nums: List[int]) -> bool:
        # STEP 1: Initialize hash set to track seen numbers
        # - Set gives O(1) lookup; tracks uniqueness efficiently
        seen = set()

        # STEP 2: Iterate through each number
        # - If num already in set -> duplicate found -> return True
        for num in nums:
            if num in seen:
                return True
            seen.add(num) # Add to set for future checks

        # STEP 3: No duplicates found during iteration
        # - Return False only after full traversal
        return False
        # -> below mentioned code also do the trick
        # return True if len(set(nums)) != len(num) else False
```

```
# ----- INLINE TESTS -----
if __name__ == "__main__":
    sol = Solution()

    # Test 1: Normal case - duplicate exists
    assert sol.containsDuplicate([1,2,3,1]) == True, "Test 1 Failed"
    print(" Test 1 passed: [1,2,3,1] → True")

    # Test 2: Edge case - single element
    assert sol.containsDuplicate([1]) == False, "Test 2 Failed"
    print(" Test 2 passed: [1] → False")

    # Test 3: Tricky/negative - all unique, large input
    assert sol.containsDuplicate([i for i in range(1000)]) == False, \
        "Test 3 Failed"
    print(" Test 3 passed: 0..999 (all unique) → False")
```

How to use: Copy-paste this block into .py or Quarto cell → run directly → instant feedback.

Example Walkthrough

Let's walk through `nums = [1, 2, 3, 1]` step by step:

Initial State:

- `seen = set()` — empty
- `nums = [1, 2, 3, 1]`

Step 1: Process `num = 1`

- Check: Is 1 in `seen`? → No
- Action: Add 1 to `seen` → `seen = {1}`
- Why? We record that we've seen 1 so future 1s will trigger duplicate.

Step 2: Process `num = 2`

- Check: Is `2` in `seen`? → No
- Action: Add `2` → `seen = {1, 2}`
- Why? Still building our uniqueness tracker.

Step 3: Process `num = 3`

- Check: Is `3` in `seen`? → No
- Action: Add `3` → `seen = {1, 2, 3}`
- Why? No duplicate yet — continue.

Step 4: Process `num = 1` (again)

- Check: Is `1` in `seen`? → Yes!
- Action: Immediately `return True`
- Why? We found a duplicate — no need to process further.
- Pattern Insight: Hashing lets us detect reoccurrence in $O(1)$, breaking early saves time.

Final Output: `True`

This exemplifies the **Arrays & Hashing** pattern:

- Use a hash set to track what we've seen.
- Leverage $O(1)$ average lookup to detect duplicates on the fly.
- Early termination optimizes best-case performance.

Complexity Analysis

- **Time Complexity:** $O(n)$

We iterate through the array once. Each `in` check and `add` operation on the set is $O(1)$ average case. Worst-case total = $O(n)$.

- **Space Complexity:** $O(n)$

In the worst case (all elements unique), we store all `n` elements in the set. Space scales linearly with input size.

3. Majority Element

Pattern: Arrays & Hashing

Problem Statement

Given an array `nums` of size `n`, return the majority element.

The majority element is the element that appears more than $n / 2$ times. You may assume that the majority element always exists in the array.

Sample Input & Output

```
Input: nums = [3,2,3]
Output: 3
Explanation: 3 appears 2 times > 3/2 = 1 → majority.
```

```
Input: nums = [2,2,1,1,1,2,2]
Output: 2
Explanation: 2 appears 4 times > 7/2 = 3 → majority.
```

```
Input: nums = [1]
Output: 1
Explanation: Single element is always majority by definition.
```

LeetCode Editorial Solution + Inline Tests


```

from typing import List

class Solution:
    def majorityElement(self, nums: List[int]) -> int:
        # STEP 1: Initialize hash map to count occurrences
        # - Tracks frequency of each element for quick lookup
        count = {}

        # STEP 2: Iterate and count each element
        # - Invariant: after each step, count[x] = freq of x so far
        for num in nums:
            count[num] = count.get(num, 0) + 1

            # Early exit: if any count exceeds n//2, return immediately
            # - Optimization: avoids full scan if found early
            if count[num] > len(nums) // 2:
                return num

        # STEP 3: Return fallback (problem guarantees existence)
        # - Should never reach here per problem constraint
        return nums[0]

# ----- INLINE TESTS -----
if __name__ == "__main__":
    sol = Solution()

    # Test 1: Normal case
    result1 = sol.majorityElement([3, 2, 3])
    print(f"Test 1 - Input: [3,2,3] → Output: {result1} (Expected: 3)")
    assert result1 == 3, "Test 1 Failed"

    # Test 2: Edge case - single element
    result2 = sol.majorityElement([1])
    print(f"Test 2 - Input: [1] → Output: {result2} (Expected: 1)")
    assert result2 == 1, "Test 2 Failed"

    # Test 3: Tricky case - long sequence, majority late
    result3 = sol.majorityElement([1, 1, 1, 2, 2, 2, 2])
    print(f"Test 3 - Input: [1,1,1,2,2,2,2] → Output: {result3} "
          f"(Expected: 2)")
    assert result3 == 2, "Test 3 Failed"

```

```
print(" All tests passed")
```

How to use: Copy-paste this block into .py or Quarto cell → run directly → instant feedback.

Example Walkthrough

Let's walk through `nums = [2,2,1,1,1,2,2]` step by step:

Initial State:

- `count = {}` (empty dictionary)
 - `len(nums) = 7`, so majority threshold = $7//2 = 3$
-

Step 1: `num = 2`

- `count.get(2, 0) + 1 = 0 + 1 = 1`
- `count = {2: 1}`
- `1 > 3?` → No → continue

Step 2: `num = 2`

- `count[2] = 1 + 1 = 2`
- `count = {2: 2}`
- `2 > 3?` → No → continue

Step 3: `num = 1`

- `count[1] = 0 + 1 = 1`
- `count = {2:2, 1:1}`
- `1 > 3?` → No → continue

Step 4: `num = 1`

- `count[1] = 1 + 1 = 2`
- `count = {2:2, 1:2}`
- `2 > 3?` → No → continue

Step 5: `num = 1`

- `count[1] = 2 + 1 = 3`
- `count = {2:2, 1:3}`
- `3 > 3?` → No → continue (note: must be $>$, not \geq)

Step 6: `num = 2`

- `count[2] = 2 + 1 = 3`

→ count = {2:3, 1:3}
→ 3 > 3? → No → continue

Step 7: num = 2

→ count[2] = 3 + 1 = 4
→ count = {2:4, 1:3}
→ 4 > 3? → YES → **return 2**

Why each step matters: - **Counting with dict:** Lets us track exact frequencies — core of hashing pattern. - **Early return:** Minor optimization — not required, but good practice. - **No fallback needed:** Problem guarantees majority exists, so we *will* hit return in loop.

Pattern Insight:

This is classic **Arrays & Hashing** — use a hash map to trade space for time, avoiding nested loops. Instead of $O(n^2)$, we get $O(n)$ by storing what we've seen.

Complexity Analysis

- **Time Complexity:** $O(n)$

Single pass through array. Each dictionary **get** and assignment is $O(1)$ average case.

- **Space Complexity:** $O(n)$

In worst case, all elements are distinct until majority is found — hash map stores up to n keys.

(Note: Boyer-Moore algorithm can solve this in $O(1)$ space — but that's a different pattern.) -> stick with hashing for cleaner implementation.

4. Valid Anagram

Pattern: Arrays & Hashing

Problem Statement

Given two strings **s** and **t**, return **true** if **t** is an anagram of **s**, and **false** otherwise.

An **anagram** is a word or phrase formed by rearranging the letters of a different word or phrase, typically using all the original letters exactly once.

Sample Input & Output

```
Input: s = "anagram", t = "nagaram"
```

```
Output: true
```

```
Explanation: All characters in 't' are a rearrangement of 's'.
```

```
Input: s = "rat", t = "car"
```

```
Output: false
```

```
Explanation: 't' has different characters than 's'.
```

```
Input: s = "", t = ""
```

```
Output: true
```

```
Explanation: Two empty strings are trivially anagrams.
```

LeetCode Editorial Solution + Inline Tests

```
from typing import List

class Solution:
    def isAnagram(self, s: str, t: str) -> bool:
        # STEP 1: Initialize structures
        # - Use two hash maps (dicts) to count char frequencies in s and t.
        # - More flexible than array - handles any Unicode char.
        if len(s) != len(t):
            return False

        count_s, count_t = {}, {}
```

```

# STEP 2: Main loop - populate both hash maps
#   - Traverse by index since strings are same length.
#   - Use .get() to safely handle missing keys.
for i in range(len(s)):
    count_s[s[i]] = 1 + count_s.get(s[i], 0)
    count_t[t[i]] = 1 + count_t.get(t[i], 0)

# STEP 3: Update state / bookkeeping
#   - No incremental update needed - we build full maps first.
#   - Final comparison validates anagram property.

# STEP 4: Return result
#   - Direct dict equality check - Python compares keys & values.
return count_s == count_t

# ----- INLINE TESTS -----
if __name__ == "__main__":
    sol = Solution()

    # Test 1: Normal case
    assert sol.isAnagram("anagram", "nagaram") == True, "Normal case failed"
    print(" Test 1 passed: 'anagram' vs 'nagaram'")

    # Test 2: Edge case - empty strings
    assert sol.isAnagram("", "") == True, "Empty strings should be anagrams"
    print(" Test 2 passed: empty strings")

    # Test 3: Tricky/negative - different letters
    assert sol.isAnagram("rat", "car") == False, "Mismatched letters"
    print(" Test 3 passed: 'rat' vs 'car'")

```

How to use: Copy-paste this block into .py or Quarto cell → run directly → instant feedback.

Example Walkthrough

Let's walk through `s = "anagram"`, `t = "nagaram"`:

Initial State: `count_s = {}`, `count_t = {}`, and `len(s) == len(t)` → proceed.

After i=0 (s[0]='a', t[0]='n'):
→ count_s['a'] = 1 + 0 → {'a': 1}
→ count_t['n'] = 1 + 0 → {'n': 1}

After i=1 (s[1]='n', t[1]='a'):
→ count_s['n'] = 1 + 0 → {'a':1, 'n':1}
→ count_t['a'] = 1 + 0 → {'n':1, 'a':1}

Continue to end:
→ 'a' appears 3x in both → count_s['a']=3, count_t['a']=3
→ 'g', 'r', 'm', 'n' each appear 1x → all match

Final Comparison:
→ count_s == {'a':3, 'n':1, 'g':1, 'r':1, 'm':1}
→ count_t == {'n':1, 'a':3, 'g':1, 'r':1, 'm':1}
→ Dictionaries are equal → return True

Why necessary?
→ Length check avoids unnecessary work and index errors.
→ .get(key, 0) prevents KeyError — critical for robust hashing.
→ Direct dict1 == dict2 leverages Python's deep equality — clean and readable.
→ This exemplifies **frequency counting with hash maps** — a foundational Arrays & Hashing technique.

Complexity Analysis

- **Time Complexity:** $O(n)$

Single pass over both strings (n steps), then $O(1)$ dict comparison (since alphabet size is bounded — even for Unicode, in practice it's limited per input).

- **Space Complexity:** $O(k)$

Where k = number of unique characters. In worst case, $k = n$ (all chars distinct). For lowercase English, $k \leq 26 \rightarrow$ effectively $O(1)$. For general Unicode, $O(k)$ where k scales with input diversity.

5. Group Anagrams

Pattern: Arrays & Hashing

Problem Statement

Given an array of strings `strs`, group the anagrams together. You can return the answer in **any order**.

An Anagram is a word or phrase formed by rearranging the letters of a different word or phrase, typically using all the original letters exactly once.

Constraints: - $1 \leq \text{strs.length} \leq 10^4$ - $0 \leq \text{strs}[i].\text{length} \leq 100$ - `strs[i]` consists of lowercase English letters.

Sample Input & Output

Input: `strs = ["eat","tea","tan","ate","nat","bat"]`

Output: `[["eat","tea","ate"],["tan","nat"],["bat"]]`

Explanation: Words with same frequency signature grouped together.

Input: `strs = [""]`

Output: `[[""]]`

Explanation: Single empty string forms its own group.

Input: `strs = ["a"]`

Output: `[["a"]]`

Explanation: Single letter trivially grouped.

LeetCode Editorial Solution + Inline Tests

```
from typing import List
from collections import defaultdict

class Solution:
    def groupAnagrams(self, strs: List[str]) -> List[List[str]]:
        # STEP 1: Create hashmap to group by freq signature
        anagram_map = defaultdict(list) # -> defaultdict(list, {})
```

```

    for word in strs:
        # STEP 2: Count frequency of 26 letters
        freq = [0] * 26
        for ch in word:
            freq[ord(ch) - ord('a')] += 1

        # STEP 3: Use tuple of freq as key
        key = tuple(freq)
        anagram_map[key].append(word)

    # STEP 4: Return grouped anagrams
    return list(anagram_map.values())

# ----- INLINE TESTS -----
if __name__ == "__main__":
    sol = Solution()

    # Test 1: Normal case
    strs1 = ["eat", "tea", "tan", "ate", "nat", "bat"]
    print(sol.groupAnagrams(strs1))
    # Expected: groups like ["eat", "tea", "ate"], ["tan", "nat"], ["bat"]

    # Test 2: Edge case - single empty string
    strs2 = [""]
    print(sol.groupAnagrams(strs2))
    # Expected: [[""]]

    # Test 3: Tricky case - single letter
    strs3 = ["a"]
    print(sol.groupAnagrams(strs3))
    # Expected: [["a"]]

```

How to use: Copy-paste this block into .py or Quarto cell → run directly → instant feedback.

Detailed Example Walkthrough

Pick input: ["eat", "tea", "tan", "ate", "nat", "bat"].

1. **Initialize:** `anagram_map = {}` (empty defaultdict). Tracks groups keyed by letter-frequency tuple.
2. **Word = “eat”**
 - `freq = [1,0,0,0,1,0,...,1,...]`
 - `key = (1,0,0,0,1,0,...,1,...)`
 - Add \rightarrow `anagram_map[key] = ["eat"]`.
3. **Word = “tea”**
 - Same frequency tuple as "eat".
 - Append \rightarrow `["eat","tea"]`.
4. **Word = “tan”**
 - freq different \rightarrow new key.
 - Add \rightarrow `["tan"]`.
5. **Word = “ate”**
 - Same key as "eat".
 - Append \rightarrow `["eat","tea","ate"]`.
6. **Word = “nat”**
 - Same key as "tan".
 - Append \rightarrow `["tan","nat"]`.
7. **Word = “bat”**
 - New frequency key \rightarrow `["bat"]`.
8. **Final Output:** `[["eat","tea","ate"],["tan","nat"],["bat"]]`.

Pattern Insight: We reduce each string into a **hashable signature** (frequency tuple). Hashing ensures $O(1)$ group lookup, which is the core of **Arrays & Hashing** mastery.

Complexity Analysis

- **Time Complexity:**

- For each word (n words), count letters (k length).
- Total: $O(n * k)$.

- **Space Complexity:**

- Hashmap stores up to n groups.
 - Each key is a 26-length tuple.
 - Total: $O(n * k)$ for output + $O(26)$ per key.
-

6. Longest Consecutive Sequence

Pattern: Arrays & Hashing

Problem Statement

Given an unsorted array of integers `nums`, return the length of the longest consecutive elements sequence.

You must write an algorithm that runs in $O(n)$ time.

Sample Input & Output

```
Input: nums = [100, 4, 200, 1, 3, 2]
```

```
Output: 4
```

```
Explanation: The longest consecutive sequence is [1, 2, 3, 4] → length 4.
```

```
Input: nums = [0, 3, 7, 2, 5, 8, 4, 6, 0, 1]
```

```
Output: 9
```

```
Explanation: Longest sequence is [0,1,2,3,4,5,6,7,8] → length 9.
```

Input: nums = []
Output: 0
Explanation: Empty input → no sequence → length 0.

LeetCode Editorial Solution + Inline Tests

```
from typing import List

class Solution:
    def longestConsecutive(self, nums: List[int]) -> int:
        # STEP 1: Initialize structures
        # Use a set for O(1) lookups. Avoid duplicates.
        num_set = set(nums)
        max_len = 0

        # STEP 2: Main loop / recursion
        # Only start counting if 'n' is the start of a sequence
        # (i.e., n-1 not in set). Avoids recounting same seq.
        for n in num_set:
            if n - 1 not in num_set:
                curr_num = n
                curr_len = 1

                # STEP 3: Update state / bookkeeping
                # Extend sequence as long as next number exists
                while curr_num + 1 in num_set:
                    curr_num += 1
                    curr_len += 1

                # Update global max
                max_len = max(max_len, curr_len)

        # STEP 4: Return result
        # Handles edge case: empty input → max_len remains 0
        return max_len

# ----- INLINE TESTS -----
if __name__ == "__main__":
```

```

sol = Solution()

# Test 1: Normal case
result1 = sol.longestConsecutive([100, 4, 200, 1, 3, 2])
print(f"Test 1 - Expected: 4, Got: {result1}")
assert result1 == 4, "Test 1 Failed"

# Test 2: Edge case
result2 = sol.longestConsecutive([])
print(f"Test 2 - Expected: 0, Got: {result2}")
assert result2 == 0, "Test 2 Failed"

# Test 3: Tricky/negative
result3 = sol.longestConsecutive([0, -1, 1, -2, 2, -3, 3])
print(f"Test 3 - Expected: 7, Got: {result3}")
assert result3 == 7, "Test 3 Failed"

print(" All inline tests passed!")

```

How to use: Copy-paste this block into .py or Quarto cell → run directly → instant feedback.

Example Walkthrough

Let's walk through `nums = [100, 4, 200, 1, 3, 2]`.

Initial state:

- `num_set = {100, 4, 200, 1, 3, 2}`
- `max_len = 0`

Step-by-step:

1. **n = 100**
 - Check: Is 99 in set? No → start sequence.
 - `curr_num = 100, curr_len = 1`
 - Check 101? Not in set → sequence ends.
 - `max_len = max(0, 1) → 1`
2. **n = 4**
 - Is 3 in set? Yes → skip (not start of sequence).

3. **n = 200**

→ 199? No → start sequence.

→ curr_len = 1, 201? No → update max_len = max(1,1) → 1

4. **n = 1**

→ 0? No → start sequence.

→ curr_num = 1, curr_len = 1

→ 2? Yes → curr_num=2, curr_len=2

→ 3? Yes → curr_num=3, curr_len=3

→ 4? Yes → curr_num=4, curr_len=4

→ 5? No → stop.

→ max_len = max(1,4) → 4

5. **n = 3** → 2 exists → skip.

6. **n = 2** → 1 exists → skip.

Final max_len = 4.

Why skip if n-1 exists?

→ Avoids recounting the same sequence multiple times.

→ Ensures each sequence is counted only from its minimum element.

→ Critical for achieving $O(n)$ — each element visited at most twice.

Pattern Insight:

Hashing lets us check existence in $O(1)$.

By anchoring sequences at their start, we avoid nested loops → linear time.

Complexity Analysis

- **Time Complexity:** $O(n)$

Each element is visited at most twice: once in outer loop, once in while loop.

Set lookups are $O(1)$. Total operations bounded by $2n \rightarrow O(n)$.

- **Space Complexity:** $O(n)$

We store all elements in a set. No recursion or auxiliary structures beyond that. Space scales linearly with input size.

7. Subarray Sum Equals K

Pattern: Arrays & Hashing

Problem Statement

Given an array of integers `nums` and an integer `k`, return the total number of subarrays whose sum equals `k`.

A subarray is a contiguous non-empty sequence of elements within an array.

Sample Input & Output

```
Input: nums = [1, 1, 1], k = 2
```

```
Output: 2
```

```
Explanation: Subarrays [1,1] (indices 0-1) and [1,1] (indices 1-2)
both sum to 2.
```

```
Input: nums = [1, 2, 3], k = 3
```

```
Output: 2
```

```
Explanation: Subarrays [1,2] (indices 0-1) and [3] (index 2) both sum to 3.
```

```
Input: nums = [1], k = 0
```

```
Output: 0
```

```
Explanation: Only one element, no way to get sum 0.
```

LeetCode Editorial Solution + Inline Tests

```

from typing import List

class Solution:
    def subarraySum(self, nums: List[int], k: int) -> int:
        # STEP 1: Initialize structures
        # - prefix_sum: running total from start to current index
        # - count: total valid subarrays found
        # - prefix_freq: tracks how many times each prefix sum occurred
        prefix_sum = 0
        count = 0
        prefix_freq = {0: 1} # crucial: empty prefix has sum 0

        # STEP 2: Main loop / recursion
        # - For each num, update running sum
        # - Check if (current_sum - k) was seen before -> means
        #   there's a subarray ending here with sum = k
        for num in nums:
            prefix_sum += num

            # STEP 3: Update state / bookkeeping
            # - If (prefix_sum - k) exists, add its frequency to count
            target = prefix_sum - k
            if target in prefix_freq:
                count += prefix_freq[target]

            # Record current prefix_sum occurrence
            prefix_freq[prefix_sum] = (
                prefix_freq.get(prefix_sum, 0) + 1
            )

        # STEP 4: Return result
        # - All edge cases handled by initializing {0:1} and using
        #   .get() with default 0
        return count

# ----- INLINE TESTS -----
if __name__ == "__main__":
    sol = Solution()

    # Test 1: Normal case
    result1 = sol.subarraySum([1, 1, 1], 2)
    print(f"Test 1 - [1,1,1], k=2 -> {result1}") # Expected: 2

```

```
# Test 2: Edge case (no matches)
result2 = sol.subarraySum([1], 0)
print(f"Test 2 - [1], k=0 → {result2}")      # Expected: 0

# Test 3: Tricky/negative (negative numbers)
result3 = sol.subarraySum([-1, -1, 1], 0)
print(f"Test 3 - [-1,-1,1], k=0 → {result3}") # Expected: 1
```

How to use: Copy-paste this block into .py or Quarto cell → run directly → instant feedback.

Example Walkthrough

Let's trace `nums = [1, 1, 1]`, `k = 2` — step by slow step.

Initial State (Before Loop):

- `prefix_sum = 0`
- `count = 0`
- `prefix_freq = {0: 1}` ← This lets us count subarrays starting at index 0!

Step 1: Process first element → `num = 1`

- Execute: `prefix_sum += 1` → now `prefix_sum = 1`
- Calculate: `target = 1 - 2 = -1`
- Check: Is `-1` in `prefix_freq`? → No → `count` unchanged (still 0)
- Update map: `prefix_freq[1] = 0 + 1 = 1`
- Map now: `{0:1, 1:1}`

State after Step 1:

- `prefix_sum = 1`
- `count = 0`
- `prefix_freq = {0:1, 1:1}`

Step 2: Process second element \rightarrow num = 1

- Execute: `prefix_sum += 1` \rightarrow now `prefix_sum = 2`
- Calculate: `target = 2 - 2 = 0`
- Check: Is 0 in `prefix_freq`? \rightarrow Yes! Value is 1
- So: `count += 1` \rightarrow `count = 1`
- Update map: `prefix_freq[2] = 0 + 1 = 1`
- Map now: {0:1, 1:1, 2:1}

State after Step 2: - `prefix_sum = 2` - `count = 1` - `prefix_freq = {0:1, 1:1, 2:1}`

Step 3: Process third element \rightarrow num = 1

- Execute: `prefix_sum += 1` \rightarrow now `prefix_sum = 3`
- Calculate: `target = 3 - 2 = 1`
- Check: Is 1 in `prefix_freq`? \rightarrow Yes! Value is 1
- So: `count += 1` \rightarrow `count = 2`
- Update map: `prefix_freq[3] = 0 + 1 = 1`
- Map now: {0:1, 1:1, 2:1, 3:1}

Final State:

- `prefix_sum = 3`
- `count = 2`
- `prefix_freq = {0:1, 1:1, 2:1, 3:1}`

Function returns 2.

Why This Works — Pattern Insight:

Arrays & Hashing Pattern: We trade space for time. Instead of checking every subarray $O(n^2)$, we store prefix sums and “ask the past” — *“Have I seen a sum that, if subtracted from now, gives me k?”*

The key equation:

$$\begin{aligned}\text{sum}[i:j] &= \text{prefix}[j] - \text{prefix}[i] = k \\ \rightarrow \text{prefix}[i] &= \text{prefix}[j] - k\end{aligned}$$

{0:1} handles subarrays starting at index 0 — because $\text{prefix}[j] - 0 = k \rightarrow$ subarray from start to j.

Complexity Analysis

- **Time Complexity:** $O(n)$

One pass through the array. Each dictionary lookup/insertion is average $O(1)$.

- **Space Complexity:** $O(n)$

In worst case (all prefix sums unique), we store $n+1$ keys in the hashmap.

8. Contiguous Array

Pattern: Arrays & Hashing

Problem Statement

Given a binary array `nums`, return the maximum length of a contiguous subarray with an equal number of 0 and 1.

Sample Input & Output

```
Input: nums = [0,1]
Output: 2
Explanation: [0, 1] is the longest contiguous subarray with equal 0s and 1s.
```

```
Input: nums = [0,1,0]
Output: 2
Explanation: [0, 1] or [1, 0] - both length 2.
No longer valid subarray exists.
```

Input: nums = [0,0,0,0,0]

Output: 0

Explanation: No 1s → no balanced subarray possible.

LeetCode Editorial Solution + Inline Tests

```
from typing import List

class Solution:
    def findMaxLength(self, nums: List[int]) -> int:
        # STEP 1: Initialize structures
        # - Use hashmap to store first occurrence of prefix balance.
        # - Key: running balance (count_1 - count_0), Value: index.
        # - Start with balance 0 at index -1 to handle full-array case.
        balance_map = {0: -1}
        max_len = 0
        balance = 0 # tracks (count of 1s) - (count of 0s)

        # STEP 2: Main loop / recursion
        # - Traverse each element. Treat 0 as -1, 1 as +1.
        # - If same balance seen before → subarray between indices
        #   has net zero → equal 0s and 1s.
        for i in range(len(nums)):
            # Update balance: +1 for 1, -1 for 0
            balance += 1 if nums[i] == 1 else -1

            # STEP 3: Update state / bookkeeping
            # - If balance seen → calculate length from first index.
            # - Else, record first occurrence of this balance.
            if balance in balance_map:
                length = i - balance_map[balance]
                if length > max_len:
                    max_len = length
            else:
                balance_map[balance] = i

        # STEP 4: Return result
        # - Handles edge cases: all 0s, all 1s, empty → returns 0.
```

```

        return max_len

# ----- INLINE TESTS -----
if __name__ == "__main__":
    sol = Solution()

    # Test 1: Normal case
    result1 = sol.findMaxLength([0, 1, 0, 1])
    print(f"Test 1 - Input: [0,1,0,1] → Output: {result1}")
    assert result1 == 4, f"Expected 4, got {result1}"

    # Test 2: Edge case - all zeros
    result2 = sol.findMaxLength([0, 0, 0, 0])
    print(f"Test 2 - Input: [0,0,0,0] → Output: {result2}")
    assert result2 == 0, f"Expected 0, got {result2}"

    # Test 3: Tricky/negative - single element
    result3 = sol.findMaxLength([1])
    print(f"Test 3 - Input: [1] → Output: {result3}")
    assert result3 == 0, f"Expected 0, got {result3}"

    print(" All inline tests passed.")

```

How to use: Copy-paste this block into .py or Quarto cell → run directly → instant feedback.

Example Walkthrough

Let's walk through `nums = [0, 1, 0, 1]` step by step.

Initial State

- `balance_map = {0: -1}` ← We start assuming balance 0 occurred “before” index 0.
 - `max_len = 0`
 - `balance = 0`
-

Step 1: $i = 0$, $\text{nums}[0] = 0$

- $\text{balance} += -1 \rightarrow \text{balance} = -1$
 - Is -1 in balance_map ? No \rightarrow record it: $\text{balance_map}[-1] = 0$
 - \rightarrow *State*: $\text{balance_map} = \{0: -1, -1: 0\}$, $\text{max_len} = 0$, $\text{balance} = -1$
-

Step 2: $i = 1$, $\text{nums}[1] = 1$

- $\text{balance} += +1 \rightarrow \text{balance} = 0$
 - Is 0 in balance_map ? Yes \rightarrow at index -1
 - $\text{Length} = 1 - (-1) = 2 \rightarrow$ update $\text{max_len} = 2$
 - \rightarrow *State*: balance_map unchanged, $\text{max_len} = 2$, $\text{balance} = 0$
-

Step 3: $i = 2$, $\text{nums}[2] = 0$

- $\text{balance} += -1 \rightarrow \text{balance} = -1$
 - Is -1 in balance_map ? Yes \rightarrow at index 0
 - $\text{Length} = 2 - 0 = 2 \rightarrow$ not greater than max_len (still 2)
 - \rightarrow *State*: unchanged max_len , $\text{balance} = -1$
-

Step 4: $i = 3$, $\text{nums}[3] = 1$

- $\text{balance} += +1 \rightarrow \text{balance} = 0$
 - Is 0 in balance_map ? Yes \rightarrow at index -1
 - $\text{Length} = 3 - (-1) = 4 \rightarrow$ update $\text{max_len} = 4$
 - \rightarrow *State*: $\text{max_len} = 4$, $\text{balance} = 0$
-

Final Output: 4

The entire array $[0, 1, 0, 1]$ is balanced \rightarrow 2 zeros, 2 ones.

Complexity Analysis

- **Time Complexity:** $O(n)$

Single pass through the array. Hashmap lookups and inserts are $O(1)$ average.

- **Space Complexity:** $O(n)$

In worst case, we store every unique balance encountered \rightarrow up to $n+1$ entries (including initial 0).
