

Heap

Pattern: Heap / Priority Queue (Min-Heap / Max-Heap)

How to Recognize

- You're asked to find **top K elements**, **kth smallest/largest**, or **median**.
- There's a need to maintain a **running order** or **priority** among elements.
- The problem involves **frequent insertions and deletions** of extremes (min/max).
- Often paired with **sorting**, **frequency counting**, or **streaming data**.

Step-by-Step Thinking Process (Template)

1. **Identify what you want to track:** e.g., k largest, k closest, top frequent.
2. **Choose the right heap type:**
 - Min-heap: keep smallest k elements → pop when size > k
 - Max-heap: keep largest k elements → use negative values in Python (min-heap trick)
3. **Use a heap of size K** to maintain only relevant candidates.
4. **Pop or push based on comparison logic.**
5. **Extract result** after processing all inputs (e.g., return root or sort remaining).

Common Pitfalls & Edge Cases

- Forgetting that Python `heapq` is a **min-heap only** → use negative values for max-heap.
- Not limiting heap size → leads to $O(n \log n)$ instead of $O(k \log n)$.
- Incorrectly handling ties (e.g., in “Top K Frequent Words”, tie-breaking by lexicographic order).
- Empty input → handle early return.

1. K Closest Points to Origin

Problem Summary

Given an array of points in 2D space, return the k closest points to the origin (0, 0) based on Euclidean distance.

Pattern

- **Heap / Priority Queue** (max-heap of size k)
- Alternative: **Sorting** (but less efficient for large datasets)

Solution with Inline Comments

```
import heapq
from typing import List, Tuple

def kClosest(points: List[List[int]], k: int) -> List[List[int]]:
    # Use a max-heap to store the k closest points
    # We store (-distance, point) so that the farthest
    # (largest distance) is at top
    # Negative distance ensures we simulate max-heap behavior using min-heap
    heap = []

    for x, y in points:
        # Calculate squared distance (avoid sqrt for speed & precision)
        dist = x*x + y*y

        # If heap has fewer than k elements, add current point
        if len(heap) < k:
            heapq.heappush(heap, (-dist, [x, y]))
        # Else, if current point is closer than the farthest in heap, replace it
        elif dist < -heap[0][0]: # -heap[0][0] is the max distance in heap
            heapq.heappop(heap)
            heapq.heappush(heap, (-dist, [x, y]))

    # Extract points from heap (they are in no particular order)
    return [point for _, point in heap]
```

```
# ---- Official LeetCode Example ----
if __name__ == "__main__":
    # Example Input: points = [[1,3],[-2,2]], k = 1
    points = [[1, 3], [-2, 2]]
    k = 1

    # Call function
    result = kClosest(points, k)

    # Expected Output: [[-2,2]]
    # Because distance of (1,3): 1+9=10; (-2,2): 4+4=8 → (-2,2) is closer
    print("Output:", result) # Output: [[-2, 2]]
```

Example Walkthrough

- Input: `points = [[1,3],[-2,2]]`, `k = 1`
- Process (1,3): $\text{dist} = 1^2 + 3^2 = 10 \rightarrow \text{heap} = [(-10, [1,3])]$
- Process (-2,2): $\text{dist} = 4 + 4 = 8 \rightarrow 8 < 10 \rightarrow \text{pop } (-10, \dots), \text{push } (-8, [-2,2])$
- Final heap: $[(-8, [-2,2])] \rightarrow \text{return } [[-2, 2]]$

Complexity

- **Time:** $O(n \log k)$ — each insertion/removal takes $O(\log k)$, done n times
 - **Space:** $O(k)$ — heap stores at most k elements
-

2. Find Median from Data Stream

Problem Summary

Design a data structure that supports adding integers and finding the median of all added numbers dynamically.

Pattern

- **Two Heaps:** Max-heap for left half, Min-heap for right half
- Balance sizes: difference ≤ 1
- Median = top of larger heap or average of both

Solution with Inline Comments

```
import heapq

class MedianFinder:
    def __init__(self):
        # Max-heap for smaller half (store negative values)
        self.small = [] # represents left half (max-heap via negatives)
        # Min-heap for larger half
        self.large = [] # represents right half (min-heap)

    def addNum(self, num: int) -> None:
        # Push to small (max-heap) first
        heapq.heappush(self.small, -num)

        # Ensure every number in small <= every number in large
        # If top of small > top of large, swap
        if self.small and self.large and (-self.small[0]) > self.large[0]:
            val = -heapq.heappop(self.small)
            heapq.heappush(self.large, val)

        # Balance the heaps: difference should be at most 1
        if len(self.small) > len(self.large) + 1:
            val = -heapq.heappop(self.small)
            heapq.heappush(self.large, val)
        elif len(self.large) > len(self.small) + 1:
            val = heapq.heappop(self.large)
            heapq.heappush(self.small, -val)

    def findMedian(self) -> float:
        # If heaps are same size, median is average
        if len(self.small) == len(self.large):
            return (-self.small[0] + self.large[0]) / 2.0
        # Else, median is top of larger heap
        elif len(self.small) > len(self.large):
            return -self.small[0]
        else:
            return self.large[0]

# ---- Official LeetCode Example ----
if __name__ == "__main__":
```

```
# Example Usage:
mf = MedianFinder()
mf.addNum(1)
mf.addNum(2)
print("Median after [1,2]:", mf.findMedian()) # Output: 1.5

mf.addNum(3)
print("Median after [1,2,3]:", mf.findMedian()) # Output: 2.0
```

Example Walkthrough

We'll go through this sequence:

```
mf = MedianFinder()
mf.addNum(1)
mf.addNum(2)
print(mf.findMedian()) # 1.5
mf.addNum(3)
print(mf.findMedian()) # 2.0
```

Step 1: addNum(1)

- Push -1 into `small` → `small = [-1]`, `large = []`
- No need to compare since `large` is empty.
- Size check:
 - `len(small) = 1`, `len(large) = 0` → difference is 1 → acceptable.

Final state: - `small = [-1]` (i.e., contains 1) - `large = []`

Step 2: addNum(2)

- Push -2 into `small` → `small = [-2, -1]` (min-heap of negatives → top is -2 → actual value is 2)
- Now check: is `top(small) > top(large)`?
 - But `large` is still empty → skip comparison.
- Balance sizes:

- `len(small) = 2, len(large) = 0` → difference is 2 (>1), so move one element.
- Pop from small: `val = -heapq.heappop(self.small)` → pop -2, so `val = 2`
- Push 2 into large: `large = [2]`
- Now `small = [-1], large = [2]`

Final state: - `small = [-1] → {1}` - `large = [2] → {2}`

Now both heaps differ in size by only 1 → good.

Step 3: `findMedian()` → after adding [1,2]

- `len(small) == 1, len(large) == 1` → equal sizes
- `Median = (-self.small[0] + self.large[0]) / 2.0`
 - `-self.small[0] = -(-1) = 1`
 - `self.large[0] = 2`
 - `Median = (1 + 2) / 2 = 1.5`

Output: 1.5 → Correct

Step 4: `addNum(3)`

- Push -3 into small → `small = [-3, -1]` → top is -3 → value is 3
- Check: is `top(small) > top(large)`?
 - `-self.small[0] = 3, self.large[0] = 2`
 - Is `3 > 2`? Yes → need to fix!
- So:
 - Pop from small: `val = -heapq.heappop(self.small)` → pop -3 `val = 3`
 - Push 3 into large: now `large = [2, 3]` → min-heap: `[2, 3]`
 - Now `small = [-1], large = [2, 3]`
- Recheck size balance:
 - `len(small) = 1, len(large) = 2` → difference is 1 → acceptable

Final state: - `small = [-1] → {1}` - `large = [2, 3] → {2, 3}`

Step 5: findMedian() → after [1,2,3]

- `len(small) = 1, len(large) = 2` → not equal
- Since `large` has more elements → median is `large[0] = 2`

Output: 2.0 → Correct

Summary of States

| Operation | small (max-heap) | large (min-heap) | Median |
|---------------------------|------------------|------------------|--------|
| <code>addNum(1)</code> | [-1] | [] | — |
| <code>addNum(2)</code> | [-1] | [2] | — |
| <code>findMedian()</code> | [-1] | [2] | 1.5 |
| <code>addNum(3)</code> | [-1] | [2, 3] | — |
| <code>findMedian()</code> | [-1] | [2, 3] | 2.0 |

Complexity

- **addNum:** $O(\log n)$ — heap operations
 - **findMedian:** $O(1)$
 - **Space:** $O(n)$
-

3. Merge k Sorted Lists

Problem Summary

Given k linked lists, each sorted in ascending order, merge them into one sorted list.

Pattern

- **Heap / Priority Queue** (k-way merge)
- At each step, pick the smallest head from k lists

Solution with Inline Comments

```
import heapq
from typing import List, Optional

# Definition for singly-linked list.
class ListNode:
    def __init__(self, val=0, next=None):
        self.val = val
        self.next = next

def mergeKLists(lists: List[Optional[ListNode]]) -> Optional[ListNode]:
    # Create a dummy head to simplify list construction
    dummy = ListNode(0)
    current = dummy

    # Min-heap to store (value, node) pairs
    heap = []

    # Initialize heap with the first node of each non-empty list
    for lst in lists:
        if lst:
            heapq.heappush(heap, (lst.val, lst))

    # While there are nodes in the heap
    while heap:
        # Pop the smallest element
        val, node = heapq.heappop(heap)

        # Link it to the result list
        current.next = node
        current = current.next

        # If this node has a next, push it into the heap
        if node.next:
            heapq.heappush(heap, (node.next.val, node.next))

    # Return the merged list (skip dummy)
    return dummy.next

# ---- Official LeetCode Example ----
```



```

if __name__ == "__main__":
    # Example Input: lists = [[1,4,5],[1,3,4],[2,6]]
    # Build linked lists
    l1 = ListNode(1, ListNode(4, ListNode(5)))
    l2 = ListNode(1, ListNode(3, ListNode(4)))
    l3 = ListNode(2, ListNode(6))

    lists = [l1, l2, l3]

    # Call function
    merged = mergeKLists(lists)

    # Print Output: [1,1,2,3,4,4,5,6]
    result = []
    while merged:
        result.append(merged.val)
        merged = merged.next
    print("Output:", result) # Output: [1, 1, 2, 3, 4, 4, 5, 6]

```

Let's walk through your **LeetCode “Merge k Sorted Lists”** code step by step, explaining how it works and why each part is important.

Problem Overview

You are given k sorted linked lists and need to merge them into one sorted linked list.

Example Input:

```

lists = [
    [1, 4, 5],
    [1, 3, 4],
    [2, 6]
]

```

Expected Output:

```
[1, 1, 2, 3, 4, 4, 5, 6]
```

Code Walkthrough

1. Define the ListNode Class

```
class ListNode:
    def __init__(self, val=0, next=None):
        self.val = val
        self.next = next
```

- This defines a node in a singly-linked list.
- Each node has a `val` (value) and a `next` pointer to the next node.
- If `next` is `None`, it's the end of the list.

Example: `ListNode(1, ListNode(4, ListNode(5)))` creates:
`1 → 4 → 5 → None`

2. Main Function: `mergeKLists`

```
def mergeKLists(lists: List[Optional[ListNode]]) -> Optional[ListNode]:
```

- Takes a list of `k` linked lists (`lists`) as input.
 - Returns a single merged sorted linked list.
-

3. Create Dummy Head

```
dummy = ListNode(0)
current = dummy
```

- A **dummy head** is used to simplify the logic of building the result list.
- **dummy** acts as a placeholder so we don't have to handle edge cases like empty result or setting the first node specially.
- **current** points to the current end of the result list — we'll append new nodes here.

Why? Because we can always return **dummy.next**, which skips the dummy node.

4. Initialize Min-Heap

```
heap = []
```

- We'll use a **min-heap** to efficiently get the smallest value among the current heads of all non-empty lists.

Heap stores tuples: (value, node)

Python's **heapq** compares the first element of the tuple, then second if needed.

5. Push First Nodes into Heap

```
for lst in lists:
    if lst:
        heapq.heappush(heap, (lst.val, lst))
```

- Loop over each list in **lists**.
- If the list is not empty (**lst** exists), push its **first node** (**lst**) into the heap with its value.
- So after this loop:
 - Heap contains: (1, l1_head), (1, l2_head), (2, l3_head)
 - The smallest values are at the top.

Important: We store the **node**, not just the value, so we can later access `node.next`.

6. Main Loop: While Heap Is Not Empty

```
while heap:
    val, node = heapq.heappop(heap)
```

- Extract the smallest element from the heap (based on `val`).
 - Now `node` is the next node to add to our result list.
-

7. Link Node to Result List

```
current.next = node
current = current.next
```

- Attach the selected node to the end of the result list.
- Move `current` forward to point to this new node.

After this, `current` tracks where the next node should be added.

8. Add Next Node from Same List (if exists)

```
if node.next:
    heapq.heappush(heap, (node.next.val, node.next))
```

- If the current node has a `next`, push that next node into the heap.
- This ensures we keep track of the next available node from the same list.

This maintains the invariant: heap always holds the **current head** of each non-empty list.

9. Return Merged List

```
return dummy.next
```

- Return the real start of the merged list (skip the dummy node).
-

Step-by-Step Execution (Example)

Given:

```
l1 = 1 → 4 → 5  
l2 = 1 → 3 → 4  
l3 = 2 → 6
```

Initial State:

- heap = [(1, l1), (1, l2), (2, l3)]
 - dummy → ?, current → dummy
-

Iteration 1:

- Pop (1, l1) → add 1 to result.
 - l1.next = 4, so push (4, l1.next) into heap.
 - Now heap: [(1, l2), (2, l3), (4, l1.next)]
 - Result: dummy → 1
-

Iteration 2:

- Pop (1, l2) → add 1 to result.
 - l2.next = 3, push (3, l2.next) into heap.
 - Heap: [(2, l3), (3, l2.next), (4, l1.next)]
 - Result: dummy → 1 → 1
-

Iteration 3:

- Pop (2, 13) → add 13 to result.
 - 13.next = 6, push (6, 13.next) into heap.
 - Heap: [(3, 12.next), (4, 11.next), (6, 13.next)]
 - Result: ... → 2
-

Iteration 4:

- Pop (3, 12.next) → add it.
 - 12.next.next = 4, push (4, 12.next.next)
 - Heap: [(4, 11.next), (4, 12.next.next), (6, 13.next)]
 - Result: ... → 3
-

Iteration 5:

- Pop (4, 11.next) → add it.
 - 11.next.next = 5, push (5, 11.next.next)
 - Heap: [(4, 12.next.next), (5, 11.next.next), (6, 13.next)]
 - Result: ... → 4
-

Iteration 6:

- Pop (4, 12.next.next) → add it.
 - No more next, so nothing pushed.
 - Heap: [(5, 11.next.next), (6, 13.next)]
 - Result: ... → 4
-

Iteration 7:

- Pop (5, 11.next.next) → add it.
 - No next, so no push.
 - Heap: [(6, 13.next)]
 - Result: ... → 5
-

Iteration 8:

- Pop (6, 13.next) → add it.
- Done!
- Heap is empty.

Final Result:

1 → 1 → 2 → 3 → 4 → 4 → 5 → 6

Correct!

Key Concepts & Optimizations

| Feature | Purpose |
|--------------------------|--|
| Min-Heap | Efficiently find the smallest current head across k lists in $O(\log k)$ time |
| Dummy Node | Avoid special handling for first node |
| Heap Stores Nodes | Allows us to access .next and continue traversal |
| Time Complexity | $O(N \log k)$, where N = total number of nodes, k = number of lists |
| Space Complexity | $O(k)$ for the heap |

Final Output

```
Output: [1, 1, 2, 3, 4, 4, 5, 6]
```

This matches the expected result perfectly.

Summary

Your implementation is **correct, efficient, and well-structured!** It uses the standard optimal approach for merging k sorted linked lists:

- Use a min-heap to always pick the smallest available node
- Advance only one list at a time
- Build result incrementally using a dummy head

Great job!

Let me know if you'd like a version with comments, visualization, or comparison to other approaches (like divide-and-conquer).

Complexity

- **Time:** $O(N \log k)$, where N = total nodes, k = number of lists
 - **Space:** $O(k)$ — heap holds at most k nodes
-

4. Task Scheduler

The “**Task Scheduler**” problem on LeetCode (Problem #621) is a popular algorithm challenge that tests greedy thinking and efficient use of data structures. Below is a detailed editorial-style explanation of the optimal solution.

Problem: Task Scheduler (LeetCode #621)

Problem Statement:

You are given a character array `tasks` representing the tasks a CPU needs to do, where each letter represents a different task. Tasks could be done in any order, but each task must be done in an interval of **at least n units** (where n is the cooling interval). Return the **least number of units of time** the CPU will take to finish all the tasks.

Example:

Input: `tasks = ["A","A","A","B","B","B"], n = 2`

Output: 8

Explanation:

A -> B -> idle -> A -> B -> idle -> A -> B

Key Observations

1. The only constraint is that **the same task cannot be scheduled within n units of time**.
 2. We can insert **idle cycles** or **other tasks** between two same tasks to satisfy the cooldown.
 3. The goal is to **minimize total time**, so we want to **minimize idle time**.
-

Strategy: Greedy with Math Insight

We can solve this using a **greedy approach**:

1. **Count frequency** of each task.
2. Find the task with **maximum frequency**, say `maxFreq`.
3. There will be $(\text{maxFreq} - 1)$ full **blocks** of size $(n + 1)$ between the most frequent task.

- For example, if `maxFreq = 3` and `n = 2`, we have:

A _ _ A _ _ A

→ 2 gaps of size 2 → total block size = 3 per segment → 2 full segments of length $(n+1) = 3$

4. Then, we **fill other tasks** into these gaps.
5. Finally, **add the remaining tasks** that occur with the same max frequency.

Formula

Let: - maxFreq = maximum frequency among tasks - countMaxFreq = number of tasks that have frequency = maxFreq

Then, the **minimum time** required is:

```
result = max(
    (maxFreq - 1) * (n + 1) + countMaxFreq,
    total_tasks
)
```

Why this formula?

- $(\text{maxFreq} - 1) * (n + 1) \rightarrow$ number of slots taken by the most frequent task and its required cooldown.
- $+ \text{countMaxFreq} \rightarrow$ the last occurrence of each task with max frequency comes after the last gap.
- But if there are **so many different tasks** that they fill all idle slots, the total time is just $\text{len}(\text{tasks})$.

So we take the **maximum** of the two to ensure we don't underestimate.

Example Walkthrough

Input: `tasks = ["A","A","A","B","B","B"], n = 2`

- Frequencies: $A \rightarrow 3, B \rightarrow 3$
- $\text{maxFreq} = 3, \text{countMaxFreq} = 2$ (both A and B appear 3 times)
- Total tasks = 6

Compute:

$$(3 - 1) * (2 + 1) + 2 = 2 * 3 + 2 = 6 + 2 = 8$$

Compare with total tasks: $\max(8, 6) = 8$

Output: 8

Code (Python)

```
from collections import Counter

class Solution:
    def leastInterval(self, tasks: List[str], n: int) -> int:
        freq = Counter(tasks)
        maxFreq = max(freq.values())
        countMaxFreq = sum(1 for f in freq.values() if f == maxFreq)

        result = (maxFreq - 1) * (n + 1) + countMaxFreq
        return max(result, len(tasks))
```

Absolutely! Let's walk through the **mathematical solution** to the **Task Scheduler (LeetCode #621)** with a **detailed, step-by-step example** to build deep intuition.

Problem Recap

Goal: Find the **minimum time** to execute all tasks, given that: - Each task is a letter (e.g., 'A', 'B'). - The same task must wait at least **n** units between executions. - You can interleave other tasks or insert **idle** slots.

Mathematical Formula (Recap)

The answer is:

```
result = max(  
    (maxFreq - 1) * (n + 1) + countMaxFreq,  
    total_tasks  
)
```

Where: - `maxFreq` = highest frequency of any task - `countMaxFreq` = number of tasks that have frequency = `maxFreq` - `total_tasks` = length of the `tasks` array

Example 1: Classic Case

Input:

`tasks = ["A","A","A","B","B","B"]`

`n = 2`

Step 1: Count Frequencies

| Task | Count |
|------|-------|
| A | 3 |
| B | 3 |

So: - `maxFreq` = 3 \rightarrow A and B both appear 3 times - `countMaxFreq` = 2 \rightarrow two tasks (A and B) have max frequency - `total_tasks` = 6

Step 2: Apply the Formula

We compute the **minimum required time due to the most frequent tasks**:

```

(maxFreq - 1) * (n + 1) + countMaxFreq
= (3 - 1) * (2 + 1) + 2
= 2 * 3 + 2
= 6 + 2
= 8

```

Now compare with total number of tasks:

```
max(8, 6) = 8
```

Answer = 8

Visual Schedule

We schedule the most frequent tasks first (A and B), spaced at least $n=2$ apart.

We start with A:

A _ _ A _ _ A

Each A is 3 units apart \rightarrow satisfies $n=2$ cooldown.

Now fill in B into the gaps:

A B _ A B _ A B

We used both B's in the gaps. Final B at the end.

No idle time left? Actually, we had to insert idle initially, but B filled most slots.

But wait — we still have:

```

Time:  1   2   3   4   5   6   7   8
       [A] [B] [idle] [A] [B] [idle] [A] [B]

```

Wait — actually, we can do better:

1: A
2: B
3: idle
4: A
5: B
6: idle
7: A
8: B

Yes \rightarrow 8 units.

We **can't do better** because A needs to run 3 times with 2 gaps of at least 2 \rightarrow forces structure.

So **8** is optimal.

Example 2: Many Unique Tasks (No Idle Needed)

Input:

tasks = ["A","A","B","B","C","C"]

n = 2

- Frequencies: A:2, B:2, C:2
- maxFreq = 2
- countMaxFreq = 3 (A, B, C all appear 2 times)
- total_tasks = 6

Apply formula:

$$(2 - 1) * (2 + 1) + 3 = 1 * 3 + 3 = 6$$
$$\max(6, 6) = 6$$

Answer = 6

Schedule:

We can do:

A B C A B C

Each task appears twice, separated by 2 other tasks \rightarrow satisfies $n=2$.

No idle needed! The **gaps are fully filled** by other tasks.

So total time = 6.

This shows: when other tasks can **fill the gaps**, we don't need idle time.

Example 3: One Very Frequent Task

Input:

```
tasks = ["A","A","A","A","A","A","B","C","D","E"]
```

```
n = 2
```

- A:6, B:1, C:1, D:1, E:1
- maxFreq = 6 (A)
- countMaxFreq = 1 (only A has frequency 6)
- total_tasks = 10

Compute:

```
(6 - 1) * (2 + 1) + 1 = 5 * 3 + 1 = 15 + 1 = 16  
max(16, 10) = 16
```

Answer = 16

Why 16?

We must schedule A like this:

A _ _ A _ _ A _ _ A _ _ A _ _ A

- 5 gaps between 6 A's
- Each gap is 3 units long ($n+1 = 3$)
- Total so far: $5*3 + 1$ (last A) \rightarrow but wait, structure is:

[A _ _] [A _ _] [A _ _] [A _ _] [A _ _] [A]

\rightarrow 5 full blocks of size 3 \rightarrow 15, plus last A \rightarrow but it's already included.

Actually, total length = $(\text{maxFreq} - 1) * (n + 1) + 1 = 5*3 + 1 = 16$

Now, we can fill B, C, D, E into the gaps.

Each gap has 2 empty slots $\rightarrow 5 \text{ gaps} \times 2 = 10$ idle spots.

We only have 5 other tasks \rightarrow we can place them, but still have idle time.

So minimum time is **16**, even though only 10 tasks.

We **can't compress** A's schedule.

Example 4: $n = 0$ (No Cooldown)

Input:

tasks = ["A","A","B","B"], $n = 0$

- maxFreq = 2, countMaxFreq = 2, total_tasks = 4

Formula:

```
(2 - 1)*(0 + 1) + 2 = 1*1 + 2 = 3
max(3, 4) = 4
```

Answer = 4

Schedule: A B A B → no idle needed, n=0 means no wait.

So total time = 4 → matches.

This shows why we take $\max(\dots, \text{total_tasks})$ — you can't finish faster than the number of tasks.

Example 5: All Same Task

Input: tasks = ["A","A","A"], n = 1

- maxFreq = 3, countMaxFreq = 1, total_tasks = 3

```
(3 - 1)*(1 + 1) + 1 = 2*2 + 1 = 5  
max(5, 3) = 5
```

Answer = 5

Schedule:

A _ A _ A

Time: 1 (A), 2 (idle), 3 (A), 4 (idle), 5 (A)

→ 5 units.

Perfect.

| Term | Meaning |
|------|---------|
|------|---------|

Key Insight Summary

| Term | Meaning |
|--|---|
| $(\text{maxFreq} - 1)$ | Number of full gaps between the most frequent task |
| $(n + 1)$ | Each gap + task needs at least n+1 time (1 for task, n for cooldown) |
| $+ \text{countMaxFreq}$ | The last occurrence of each max-frequency task goes at the end |
| $\text{max}(\dots, \text{total_tasks})$ | Can't take less time than total number of tasks |

When Does Each Term Win?

| Case | Which term wins? |
|----------------------------|---|
| Few unique tasks, high n | $(\text{maxFreq}-1)*(n+1)+\dots \rightarrow$ idle-heavy |
| Many unique tasks | total_tasks \rightarrow no idle, fully packed |
| $n = 0$ | Always total_tasks |
| One dominant task | Formula term wins |

Final Thoughts

The **mathematical solution** works because:

The **most frequent task** dictates the **minimum skeleton** of the schedule.

Other tasks either **fill the gaps** or extend the time.

If they fill everything, total tasks wins. Otherwise, the structure wins.

You don't need to simulate — just compute the **tightest possible schedule** around the bottleneck task.

Let me know if you'd like a **diagram version** or a **whiteboard-style explanation**!

Complexity Analysis

- **Time:** $O(m)$, where $m = \text{len(tasks)}$ — for counting frequencies.
 - **Space:** $O(1)$ — since there are at most 26 unique tasks (uppercase letters), the `Counter` uses constant space.
-

Why Greedy Works?

We always want to **schedule the most frequent tasks first** to avoid having them pile up at the end, which would require more idle time. By structuring around the most frequent task, we create a framework that minimizes idle slots.

Even if we interleave other tasks optimally, the **bottleneck** is the most frequent task and its cooldown requirement.

Edge Cases

- $n = 0$: No cooldown \rightarrow answer is len(tasks)
 - All tasks are the same: `["A","A"]`, $n=1 \rightarrow A \text{ idle } A \rightarrow 3$
 - Many distinct tasks: If $n=1$, `tasks = [A,B,C,D,E]`, then no idle needed \rightarrow answer = 5
-

Conclusion

The key insight is **not simulating** the schedule, but using **mathematical reasoning** to compute the minimum required time.

$$\text{Answer} = \max((\text{maxFreq} - 1) * (n + 1) + \text{countMaxFreq}, \text{total_tasks})$$

This elegant solution runs in linear time and is optimal.

Let me know if you'd like a **simulation-based approach** (using heap or queue) as an alternative method!

5. Top K Frequent Words

Problem :

We want to find the **k** most frequent words in a list, with ties broken by **lexicographical (dictionary) order**.

```
words = ["the","day","is","sunny","the","the","the","sunny","is","is"]
k = 4
print(topKFrequent(words, k))
# Output: ["the", "is", "sunny", "day"]
```

Solution (Python):

```
import heapq
from collections import Counter

def topKFrequent(words, k):
    # Step 1: Count frequency of each word
    count = Counter(words)

    # Step 2: Create a min-heap (or use negative frequency for max behavior)
```

```

    heap = [(-freq, word) for word, freq in count.items()]
    heapq.heapify(heap)

    # Step 3: Pop the top k elements
    return [heapq.heappop(heap)[1] for _ in range(k)]

# ---- Official LeetCode Example ----
if __name__ == "__main__":
    # Example Input: words = ["i","love","leetcode","i","love","coding"],k = 2
    words = ["i", "love", "leetcode", "i", "love", "coding"]
    k = 2

    # Call function
    result = topKFrequent(words, k)

    # Expected Output: ["i","love"]
    # i:2, love:2, coding:1 → top 2 → i and
    # love (tie broken by lex order: i < love)
    print("Output:", result) # Output: ['i', 'love']

```

Step-by-Step Breakdown

Step 1: Count Frequencies

```
count = Counter(words)
```

- Counter is a subclass of dict that counts occurrences.
- Example:

```

words = ["i","love","leetcode","i","love","coding"]
count = Counter(words)
# count = {'i': 2, 'love': 2, 'leetcode': 1, 'coding': 1}

```

Step 2: Build a List for the Heap

```
heap = [(-freq, word) for word, freq in count.items()]
```

- We create a list of tuples: `(-frequency, word)`
- We use **negative frequency** because:
 - Python's `heapq` is a **min-heap** by default.
 - To simulate a **max-heap** for frequency, we negate the frequency.
 - So higher actual frequency becomes more negative → smaller in min-heap → comes out first.

Example:

From our `count`, this gives:

```
[(-2, 'i'), (-2, 'love'), (-1, 'leetcode'), (-1, 'coding')]
```

Now, when we heapify, the **smallest tuple** (by first element, then second) will be at the top.

But here's the **key insight**:

When two frequencies are the same (e.g., -2), Python compares the **second element**, which is the word.

So: `(-2, 'i')` vs `(-2, 'love')` → `'i' < 'love'` lexicographically → `(-2, 'i')` is smaller.

- But we want **higher frequency first**, and **lexicographically smaller word first** in the result.

Wait — doesn't that mean `'i'` should come before `'love'`? Yes.

But in the heap, since `(-2, 'i')` is smaller than `(-2, 'love')`, it will be **popped first** — which is exactly what we want.

So the tuple `(-freq, word)` naturally gives us: - Higher frequency first (because of `-freq`)

- Lexicographically smaller word first in case of tie

This is why the tuple ordering works perfectly.

Step 3: Heapify the List

```
heapq.heapify(heap)
```

- Converts the list into a **min-heap** in-place.
- The smallest element (i.e., highest frequency, then lexicographically smallest) is at the top.

After heapify, the internal structure maintains heap property: - `heap[0]` is always the smallest (i.e., the “best” candidate).

But note: The entire list is **not sorted** — just heap-ordered.

Step 4: Extract Top k Elements

```
return [heapq.heappop(heap)[1] for _ in range(k)]
```

- We pop `k` times.
- Each `heappop()` removes and returns the smallest (i.e., most frequent, or lexicographically smaller) element.
- We take `[1]` → the **word** part of the tuple `(-freq, word)`.

Each pop takes $O(\log n)$ time, so `k` pops → $O(k \log n)$.

Example Walkthrough

Let’s run through:

```
words = ["i","love","leetcode","i","love","coding"]  
k = 2
```

Step 1: Count

```
count = {'i': 2, 'love': 2, 'leetcode': 1, 'coding': 1}
```

Step 2: Build heap list

```
heap = [(-2, 'i'), (-2, 'love'), (-1, 'leetcode'), (-1, 'coding')]
```

Step 3: heapify

After `heapify`, the heap is reordered so that: `(-2, 'i')` is at the top (smallest), because `'i' < 'love'` lexicographically.

So the heap order ensures: 1. `(-2, 'i')` 2. `(-2, 'love')` 3. `(-1, ...)` etc.

Step 4: Pop k=2 times

- 1st pop: `(-2, 'i') → append 'i'`
- 2nd pop: `(-2, 'love') → append 'love'`

Result: `["i", "love"]`

Why Doesn't Lex Order Mess It Up?

Suppose we had:

```
words = ["love", "i", "i", "love"] # same frequencies
```

Then: `(-2, 'i')` and `(-2, 'love')` - `'i' < 'love' → so (-2, 'i')` is smaller `→ popped first → correct.`

So the natural tuple comparison handles the tie-break correctly.

Time & Space Complexity

| Aspect | Complexity |
|--------|--|
| Time | $O(n + k \log n)$ - $O(n)$ for counting - $O(n)$ for heapify - $O(k \log n)$ for popping k times |
| Space | $O(n)$ for counter and heap |

Final Notes

- The **tuple** `(-freq, word)` is the magic key.
 - Python's **lexicographic comparison of strings in tuples** makes tie-breaking automatic.
 - `heapq` only supports min-heap, so we **negate frequency** to simulate max behavior.
-

Let me know if you'd like to see the **bucket sort** version ($O(n)$ time) too!

6. Find K Closest Elements

Problem Summary

Given a sorted array and integer **k**, return the **k** closest elements to a target value **x**. Return them in ascending order.

Pattern

- **Binary Search on Answer** (find left boundary of result window)
- **Two Pointers** (after finding start, expand outward)
- Or: **Sliding Window** on sorted array

Solution with Inline Comments

```
from typing import List

def findClosestElements(arr: List[int], k: int, x: int) -> List[int]:
    # Use binary search to find the leftmost starting index of k elements
    left, right = 0, len(arr) - k # right is len-k because we need k elements

    while left < right:
        mid = (left + right) // 2

        # Compare the distances from mid and mid+k to x
        # If arr[mid] is farther than arr[mid+k],
        # then mid cannot be the left bound
        # Because we'd get better elements by moving right
```

```

        if x - arr[mid] > arr[mid + k] - x:
            left = mid + 1
        else:
            right = mid

    # Now left is the starting index of the k closest elements
    return arr[left:left + k]

# ---- Official LeetCode Example ----
if __name__ == "__main__":
    # Example Input: arr = [1,2,3,4,5], k = 4, x = 3
    arr = [1, 2, 3, 4, 5]
    k = 4
    x = 3

    # Call function
    result = findClosestElements(arr, k, x)

    # Expected Output: [1,2,3,4]
    # Distances: |1-3|=2, |2-3|=1, |3-3|=0, |4-3|=1, |5-3|=2
    # Closest 4: 2,3,4,2 → but 1,2,3,4 are closer than 5
    print("Output:", result) # Output: [1, 2, 3, 4]

```

Example Walkthrough

Example Input

```

arr = [1, 2, 3, 4, 5]
k = 4
x = 3

```

We want the 4 closest elements to 3.

Step-by-Step Walkthrough

Step 1: Initial Setup

```
left = 0
right = len(arr) - k = 5 - 4 = 1
```

So our binary search range is $[0, 1)$ → only possible values for **left** are 0 or 1.

We are trying to find the **starting index** of a subarray of length $k=4$ that contains the closest elements to $x=3$.

Possible windows: - Start at 0 → $[1,2,3,4]$ - Start at 1 → $[2,3,4,5]$

We'll use binary search to pick the best one.

Binary Search Loop

Iteration 1:

```
left = 0, right = 1
mid = (0 + 1) // 2 = 0
```

Now compare: - $x - \text{arr}[\text{mid}]$ → distance from x to left end of window - $\text{arr}[\text{mid} + k] - x$ → distance from x to right end of window

Why this comparison?

Because we're comparing two overlapping windows: - One starting at $\text{mid} = 0$: $[1,2,3,4]$ - One starting at $\text{mid} + 1 = 1$: $[2,3,4,5]$

We decide which one is better by comparing the **outer edges**: $\text{arr}[\text{mid}]$ vs $\text{arr}[\text{mid} + k]$.

If $\text{arr}[\text{mid} + k]$ is closer to x , then we should move the window right → discard current mid .

Let's compute:

```
x - arr[mid] = 3 - arr[0] = 3 - 1 = 2
arr[mid + k] - x = arr[0 + 4] - 3 = arr[4] - 3 = 5 - 3 = 2
```

So:

```
if 2 > 2 → False
```

So we go to **else**:

```
right = mid = 0
```

Now $\text{left} = 0, \text{right} = 0 \rightarrow$ loop ends.

Final Result

```
return arr[left : left + k] = arr[0:4] = [1, 2, 3, 4]
```

Why [1,2,3,4] and not [2,3,4,5]?

Let's compute distances to $x = 3$:

| Element | Distance |
|---------|----------|
| 1 | |
| 2 | |
| 3 | |
| 4 | |
| 5 | |

Top 4 smallest distances: all except one of the 2s.

But both 1 and 5 are equally distant from 3. Since $1 < 5$, we prefer 1. So we pick [1,2,3,4].

This matches our result.

Key Insight of the Algorithm

Instead of comparing individual elements, we compare **candidate windows** of size k .

At each mid , we consider: - Window starting at mid : includes $arr[mid]$ to $arr[mid + k - 1]$ - The next window would start at $mid + 1$

To decide whether to move right, we compare: - $x - arr[mid] \rightarrow$ how far the **leftmost element** of current window is from x - $arr[mid + k] - x \rightarrow$ how far the **next element after the window** is from x

If the next element ($arr[mid+k]$) is **closer** than the current leftmost ($arr[mid]$), we should shift the window right.

Hence:

```
if x - arr[mid] > arr[mid + k] - x:
    left = mid + 1    # shift window right
else:
    right = mid       # keep current left or go left
```

Complexity

- **Time:** $O(\log(n - k))$ — binary search over $n - k$ positions
- **Space:** $O(1)$ — only indices used

7. Kth Largest Element in an Array

Problem

Given an array `nums` and integer k , find the **kth largest element**.

Example:

`nums = [3,2,1,5,6,4]`, $k = 2 \rightarrow$ return 5 (since 5 is the 2nd largest)

Why Use a Min-Heap?

We want the **kth largest**, so we only need to keep track of the **top k largest elements**.

Code

```
import heapq

class Solution:
    def findKthLargest(self, nums: list[int], k: int) -> int:
        # Min-heap to store the k largest elements
        heap = []

        for num in nums:
            if len(heap) < k:
                # If we have space, add the number
                heapq.heappush(heap, num)
            elif num > heap[0]:
                # If current number is bigger than the smallest in heap,
                # replace the smallest with this one
                heapq.heapreplace(heap, num)

        # The root of the min-heap is the kth largest
        return heap[0]
```

Step-by-Step Walkthrough with `nums = [3,2,1,5,6,4]`, `k = 2`

```
heap = [] # min-heap
```

1. `num = 3`

- `len(heap) = 0 < 2` → push 3
- `heap = [3]`

2. `num = 2`

- `len(heap) = 1 < 2` → push 2
- `heap = [2, 3]` (heap property: min at front)

3. `num = 1`

- `len(heap) = 2` → not less than k
- Is `1 > heap[0]`? → `1 > 2`? No → skip

4. **num = 5**

- $\text{len}(\text{heap}) = 2 \rightarrow \text{check if } 5 > 2 \rightarrow \text{Yes}$
- Replace: `heapreplace(heap, 5)` \rightarrow removes 2, adds 5
- $\text{heap} = [3, 5] \rightarrow$ now min is 3

5. **num = 6**

- $6 > 3 \rightarrow \text{Yes}$
- `heapreplace(heap, 6)` \rightarrow removes 3, adds 6
- $\text{heap} = [5, 6] \rightarrow$ min is 5

6. **num = 4**

- $4 > 5?$ No \rightarrow skip

Final heap: $[5, 6] \rightarrow \text{heap}[0] = 5 \rightarrow$ return **5**

Time & Space Complexity

| Metric | Complexity | Explanation |
|--------------|---------------|--|
| Time | $O(n \log k)$ | For each of n elements: heap operation takes $O(\log k)$ |
| Space | $O(k)$ | Heap stores at most k elements |

Efficient when **k is small** compared to n (e.g., $k = 10$, $n = 10000$)

Pro Tips

- Use `heapq.heapreplace()` instead of `heappop()` + `heappush()` for efficiency.
- Always compare `num > heap[0]` — not `>=`, because duplicates are allowed.
- This method works even if there are duplicate values.

Example: `nums = [1,1,1,2,2]`, $k = 3 \rightarrow$ 3rd largest is 1 \rightarrow correct.

8. Smallest Range Covering Elements from K Lists

Problem Statement:

You are given k sorted integer arrays. You need to find the **smallest range** that includes **at least one number from each array**.

The range is defined as $[\text{start}, \text{end}]$, and its **size** is $\text{end} - \text{start}$.

Return the **smallest such range**. If multiple ranges have the same size, return any one of them.

Example:

```
Input: nums = [[4,10,15,24,26], [0,9,12,20], [5,18,22,30]]
Output: [20,24]
```

Explanation: The range $[20, 24]$ covers: - 20 from the second list, - 24 from the first list, - 22 from the third list.

All lists are covered, and it's the smallest possible range.

Key Insight:

We want to minimize the difference ($\text{end} - \text{start}$) while ensuring that **each of the k lists contributes at least one element** in the range.

A greedy + heap approach works well here.

Python Implementation:

```
import heapq
from typing import List

class Solution:
    def smallestRange(self, nums: List[List[int]]) -> List[int]:
        # Min-heap to store (value, list_index, index_in_list)
        heap = []
        max_val = float('-inf')
```



```

# Initialize: add the first element from each list
for i in range(len(nums)):
    heapq.heappush(heap, (nums[i][0], i, 0))
    max_val = max(max_val, nums[i][0])

# Initialize result range
best_start, best_end = float('-inf'), float('inf')

while heap:
    min_val, list_idx, idx_in_list = heapq.heappop(heap)

    # Update the best range if current range is smaller
    if max_val - min_val < best_end - best_start:
        best_start, best_end = min_val, max_val

    # Move to next element in the same list
    if idx_in_list + 1 < len(nums[list_idx]):
        next_val = nums[list_idx][idx_in_list + 1]
        heapq.heappush(heap, (next_val, list_idx, idx_in_list + 1))
        max_val = max(max_val, next_val)
    else:
        # One list is exhausted; we can't form a valid range anymore
        break

return [best_start, best_end]

```

Complexity Analysis:

- **Time Complexity:**
 $O(N \log k)$, where N is the total number of elements across all lists, and k is the number of lists.
Each element is pushed and popped once from the heap ($\log k$ per operation).
- **Space Complexity:**
 $O(k)$ for the heap (stores one element per list at a time).

Why This Works:

- We always maintain one element from each list (initially), then replace the smallest one with the next in its list.
- By doing this, we ensure we never skip a potentially better range.

- The heap ensures we always process the smallest current element, which helps shrink the range.

Example walkthrough

We'll use this example:

```
nums = [
    [4, 10, 15, 24, 26], # List 0
    [0, 9, 12, 20],      # List 1
    [5, 18, 22, 30]      # List 2
]
```

Line-by-Line Walkthrough (With Visuals & Tracing)

Let's now go **step-by-step**, updating variables at every stage.

Step 1: Initialize heap and max_val

```
heap = []
max_val = float('-inf') # -∞
```

Now loop over each list ($i = 0, 1, 2$):

$i = 0$: List 0 \rightarrow element = 4

- Push (4, 0, 0) into heap
- $\text{max_val} = \max(-\infty, 4) = 4$

Heap: [(4, 0, 0)]

$i = 1$: List 1 \rightarrow element = 0

- Push (0, 1, 0) into heap
- $\text{max_val} = \max(4, 0) = 4$

Heap: [(0, 1, 0), (4, 0, 0)] \rightarrow min-heap sorted: [0, 4]

i = 2: List 2 → element = 5

- Push (5, 2, 0) into heap
- $\text{max_val} = \max(4, 5) = 5$

Heap: [(0, 1, 0), (4, 0, 0), (5, 2, 0)] → sorted by value

After initialization: - heap = [(0, 1, 0), (4, 0, 0), (5, 2, 0)] - $\text{max_val} = 5$ - $\text{best_start} = -\infty$, $\text{best_end} = \infty$

This window: {0 (list1), 4 (list0), 5 (list2)} → covers all lists!

Step 2: Set best_start, best_end

```
best_start, best_end = float('-inf'), float('inf')
```

So far, no valid range → we'll update it when we find a better one.

Step 3: Start the while heap: Loop

We process the heap until it's empty or a list runs out.

Let's trace each iteration.

Iteration 1: Pop (0, 1, 0)

```
min_val, list_idx, idx_in_list = heapq.heappop(heap)
# → min_val = 0, list_idx = 1, idx_in_list = 0
```

Now check:

```
if max_val - min_val < best_end - best_start:
    # 5 - 0 = 5 < ∞ - (-∞) → True
    best_start, best_end = 0, 5
```

Update best range: [0, 5] (size = 5)

Now try to advance list 1:

```
if idx_in_list + 1 < len(nums[1]): # 0+1=1 < 4 → True
    next_val = nums[1][1] = 9
    heapq.heappush(heap, (9, 1, 1))
    max_val = max(5, 9) = 9
```

New heap: [(4, 0, 0), (5, 2, 0), (9, 1, 1)]

→ Sorted: [4, 5, 9]

Now window: {4, 5, 9} → min=4, max=9 → range=5

Iteration 2: Pop (4, 0, 0)

```
min_val = 4, list_idx = 0, idx_in_list = 0
```

Check:

```
if 9 - 4 = 5 < 5 - 0 = 5? → No (5 < 5 is False)
```

No update.

Advance list 0:

```
if 0+1=1 < 5 → True
next_val = nums[0][1] = 10
push (10, 0, 1)
max_val = max(9, 10) = 10
```

Heap: [(5, 2, 0), (9, 1, 1), (10, 0, 1)] → sorted: [5, 9, 10]

Window: {5, 9, 10} → range = 5

Iteration 3: Pop (5, 2, 0)

```
min_val = 5, list_idx = 2, idx_in_list = 0
```

Check:

```
10 - 5 = 5 < 5 → False → no update
```

Advance list 2:

```
1 < 4 → True  
next_val = nums[2][1] = 18  
push(18, 2, 1)  
max_val = max(10, 18) = 18
```

Heap: [(9, 1, 1), (10, 0, 1), (18, 2, 1)] → [9, 10, 18]

Window: {9, 10, 18} → range = 9

Iteration 4: Pop (9, 1, 1)

```
min_val = 9, list_idx = 1, idx_in_list = 1
```

Check:

```
18 - 9 = 9 < 5? → No → skip
```

Advance list 1:

```
1+1=2 < 4 → True  
next_val = nums[1][2] = 12  
push(12, 1, 2)  
max_val = max(18, 12) = 18
```

Heap: [(10, 0, 1), (12, 1, 2), (18, 2, 1)] → [10, 12, 18]

Window: {10, 12, 18} → range = 8

Iteration 5: Pop (10, 0, 1)

```
min_val = 10, list_idx = 0, idx_in_list = 1
```

Check:

```
18 - 10 = 8 < 5? → No
```

Advance list 0:

```
1+1=2 < 5 → True  
next_val = nums[0][2] = 15  
push (15, 0, 2)  
max_val = max(18, 15) = 18
```

Heap: [(12, 1, 2), (15, 0, 2), (18, 2, 1)] → [12, 15, 18]

Window: {12, 15, 18} → range = 6

Iteration 6: Pop (12, 1, 2)

```
min_val = 12, list_idx = 1, idx_in_list = 2
```

Check:

```
18 - 12 = 6 < 5? → No
```

Advance list 1:

```
2+1=3 < 4 → True  
next_val = nums[1][3] = 20  
push (20, 1, 3)  
max_val = max(18, 20) = 20
```

Heap: [(15, 0, 2), (18, 2, 1), (20, 1, 3)] → [15, 18, 20]

Window: {15, 18, 20} → range = 5 → same as before → no update

Iteration 7: Pop (15, 0, 2)

```
min_val = 15, list_idx = 0, idx_in_list = 2
```

Check:

```
20 - 15 = 5 < 5? → No
```

Advance list 0:

```
2+1=3 < 5 → True  
next_val = nums[0][3] = 24  
push (24, 0, 3)  
max_val = max(20, 24) = 24
```

Heap: [(18, 2, 1), (20, 1, 3), (24, 0, 3)] → [18, 20, 24]

Window: {18, 20, 24} → range = 6

Iteration 8: Pop (18, 2, 1)

```
min_val = 18, list_idx = 2, idx_in_list = 1
```

Check:

```
24 - 18 = 6 < 5? → No
```

Advance list 2:

```
1+1=2 < 4 → True  
next_val = nums[2][2] = 22  
push (22, 2, 2)  
max_val = max(24, 22) = 24
```

Heap: [(20, 1, 3), (22, 2, 2), (24, 0, 3)] → [20, 22, 24]

Now check:

```
24 - 20 = 4 < 5? → YES!
```

Update best range: `best_start = 20, best_end = 24`

We found a better range: `[20, 24]` (size = 4)

Iteration 9: Pop (20, 1, 3)

```
min_val = 20, list_idx = 1, idx_in_list = 3
```

Check:

```
24 - 20 = 4 < 4? → No (4 == 4)
```

Now try to advance list 1:

```
3+1=4 < 4? → False → list 1 is exhausted!  
break
```

Loop ends.

Final Output

```
return [best_start, best_end] # → [20, 24]
```

| Iteration | Popped From | New Max | Current Window | Range | Best Range |
|-----------|-------------|---------|----------------|-------|------------|
|-----------|-------------|---------|----------------|-------|------------|

Summary Table: Key Variables Over Time

| Iteration | Popped From | New Max | Current Window | Range | Best Range |
|-----------|-------------|---------|------------------------|-------|------------|
| 1 | List 1 (0) | 9 | {4,5,9} | 5 | [0,5] |
| 2 | List 0 (4) | 10 | {5,9,10} | 5 | [0,5] |
| 3 | List 2 (5) | 18 | {9,10,18} | 9 | [0,5] |
| 4 | List 1 (9) | 18 | {10,12,18} | 8 | [0,5] |
| 5 | List 0 (10) | 18 | {12,15,18} | 6 | [0,5] |
| 6 | List 1 (12) | 20 | {15,18,20} | 5 | [0,5] |
| 7 | List 0 (15) | 24 | {18,20,24} | 6 | [0,5] |
| 8 | List 2 (18) | 24 | {20,22,24} | 4 | [20,24] |
| 9 | List 1 (20) | 24 | List 1 done → break | | |

Why This Works: Algorithm Logic

| Concept | Explanation |
|---------------------------------------|--|
| Min-Heap | Always picks the smallest current element → helps shrink the left side of the range. |
| Track max_val | Ensures we know how wide the current window is. |
| Replace with next in same list | Keeps one element per list, explores new combinations. |
| Break when list exhausted | Can't form a full window anymore → stop. |
| Greedy but optimal | Because arrays are sorted, advancing the smallest guarantees we don't miss the global minimum. |

Final Answer

[20, 24]

Pro Tips for Understanding

- Think of the heap as a “**priority queue**” of “front runners” — always the smallest.
- The `max_val` is like the **tallest person in the group** — we care about the span between shortest and tallest.
- Every time we move the shortest forward, we’re trying to **tighten the group**.