

Binary Search Tree

1. Validate Binary Search Tree

Pattern: Tree Traversal + Range Validation (In-Order / DFS with Bounds)

Problem Statement

Given the **root** of a binary tree, determine if it is a valid binary search tree (BST).

A **valid BST** is defined as follows:

- The left subtree of a node contains only nodes with keys **less than** the node's key.
- The right subtree of a node contains only nodes with keys **greater than** the node's key.
- Both the left and right subtrees must also be binary search trees.

Note: Duplicate values are **not allowed** in a BST per this problem.

Sample Input & Output

```
Input: root = [2,1,3]
```

```
Output: true
```

```
Explanation: 2 is root; left=1 (<2), right=3 (>2); both subtrees valid.
```

```
Input: root = [5,1,4,null,null,3,6]
Output: false
Explanation: Root=5; right child=4 (<5 ).
Also, 4's left=3 is <5 but appears in right subtree → violates BST.
```

```
Input: root = [1]
Output: true
Explanation: Single node is always a valid BST.
```

LeetCode Editorial Solution + Inline Tests

```
from typing import Optional

# Definition for a binary tree node.
class TreeNode:
    def __init__(self, val=0, left=None, right=None):
        self.val = val
        self.left = left
        self.right = right

class Solution:
    def isValidBST(self, root: Optional[TreeNode]) -> bool:
        # STEP 1: Initialize recursive helper with bounds
        # - Use -inf and +inf as initial valid range for root
        # - Each recursive call tightens the allowed range

        def validate(node, low, high):
            # STEP 2: Base case - empty node is valid
            if not node:
                return True

            # STEP 3: Check current node against bounds
            # - Must satisfy: low < node.val < high
            if node.val <= low or node.val >= high:
                return False

            # STEP 4: Recurse left and right with updated bounds
```

```

        # - Left subtree: upper bound becomes node.val
        # - Right subtree: lower bound becomes node.val
        return (validate(node.left, low, node.val) and
                validate(node.right, node.val, high))

    # STEP 5: Start validation from root with full range
    return validate(root, float('-inf'), float('inf'))

# ----- INLINE TESTS -----
if __name__ == "__main__":
    sol = Solution()

    # Test 1: Normal case - valid BST
    root1 = TreeNode(2, TreeNode(1), TreeNode(3))
    print("Test 1:", sol.isValidBST(root1)) # Expected: True

    # Test 2: Edge case - single node
    root2 = TreeNode(1)
    print("Test 2:", sol.isValidBST(root2)) # Expected: True

    # Test 3: Tricky case - invalid BST (right child too small)
    #      5
    #     / \
    #    1   4
    #     / \
    #    3   6
    root3 = TreeNode(5)
    root3.left = TreeNode(1)
    root3.right = TreeNode(4)
    root3.right.left = TreeNode(3)
    root3.right.right = TreeNode(6)
    print("Test 3:", sol.isValidBST(root3)) # Expected: False

```

How to use: Copy-paste this block into `.py` or Quarto cell → run directly → instant feedback.

Example Walkthrough

We'll walk through **Test 3** (`[5,1,4,null,null,3,6]`) step by step.

1. Call `isValidBST(root3)`

- `root3.val = 5`
- Calls `validate(root3, -inf, +inf)`

2. Inside `validate(node=5, low=-inf, high=+inf)`

- Node exists \rightarrow continue
- Check: `5 <= -inf`? No. `5 >= +inf`? No. \rightarrow OK
- Now recurse:
 - Left: `validate(1, -inf, 5)`
 - Right: `validate(4, 5, +inf)`

3. Process left subtree: `validate(1, -inf, 5)`

- 1 is between `-inf` and 5 \rightarrow OK
- Left child = `None` \rightarrow returns `True`
- Right child = `None` \rightarrow returns `True`
- Returns `True`

4. Process right subtree: `validate(4, 5, +inf)`

- Check: `4 <= 5`? **Yes** \rightarrow but condition is `node.val <= low` \rightarrow `4 <= 5` is true, but wait!
 - **Correction:** `low = 5`, so `4 <= 5` \rightarrow **true**, which triggers return `False`
 - Because in right subtree of 5, all values **must be** `> 5`, but 4 is **not** `> 5`
- So: `4 >= high`? No (`high = inf`)
But `4 <= low` (5) \rightarrow **yes** \rightarrow **invalid!**
- Returns `False`

5. **Final result:** `True` and `False` \rightarrow `False`

Output: `False` — correctly identifies invalid BST.

Complexity Analysis

- **Time Complexity:** $O(n)$

We visit every node exactly once in the worst case (skewed tree or fully valid BST).

- **Space Complexity:** $O(h)$, where h is height of tree

Due to recursion stack depth. In worst case (skewed tree), $h = n$; in balanced tree, $h = \log n$.

2. Convert Sorted Array to Binary Search Tree

Pattern: Divide and Conquer / Binary Search Tree Construction

Problem Statement

Given an integer array `nums` where the elements are sorted in **ascending order**, convert it to a **height-balanced** binary search tree.

A height-balanced binary tree is defined as a binary tree in which the depth of the two subtrees of every node never differs by more than one.

Sample Input & Output

Input: `nums = [-10, -3, 0, 5, 9]`

Output: `[0, -3, 9, -10, null, 5]`

Explanation: One possible answer is `[0, -3, 9, -10, null, 5]`, which represents the following height-balanced BST:

```
      0
     / \
    -3  9
   /  /
  -10 5
```

Input: nums = [1, 3]
Output: [3, 1] or [1, null, 3]
Explanation: Both are valid height-balanced BSTs.

Input: nums = [0]
Output: [0]
Explanation: Single-node tree is trivially balanced.

LeetCode Editorial Solution + Inline Tests

```
from typing import List, Optional

# Definition for a binary tree node.
class TreeNode:
    def __init__(self, val=0, left=None, right=None):
        self.val = val
        self.left = left
        self.right = right

class Solution:
    def sortedArrayToBST(self, nums: List[int]) -> Optional[TreeNode]:
        # STEP 1: Base case - empty subarray
        # - Return None to terminate recursion
        if not nums:
            return None

        # STEP 2: Choose middle element as root
        # - Ensures left/right subtrees differ by 1 in size
        mid = len(nums) // 2
        root = TreeNode(nums[mid])

        # STEP 3: Recursively build left and right subtrees
        # - Left: elements before mid (guaranteed < root.val)
        # - Right: elements after mid (guaranteed > root.val)
        root.left = self.sortedArrayToBST(nums[:mid])
        root.right = self.sortedArrayToBST(nums[mid + 1:])
```

```

        # STEP 4: Return constructed subtree root
        # - Base case handles empty slices automatically
        return root

# ----- INLINE TESTS -----
if __name__ == "__main__":
    sol = Solution()

    # Helper to serialize tree (preorder with None markers)
    def serialize(root):
        if not root:
            return [None]
        return [root.val] + serialize(root.left) + serialize(root.right)

    # Test 1: Normal case
    tree1 = sol.sortedArrayToBST([-10, -3, 0, 5, 9])
    ser1 = serialize(tree1)
    # Trim trailing Nones for cleaner comparison
    while ser1 and ser1[-1] is None:
        ser1.pop()
    print("Test 1:", ser1) # Expect: [0, -3, -10, None, None, None, 9, 5]

    # Test 2: Edge case - two elements
    tree2 = sol.sortedArrayToBST([1, 3])
    ser2 = serialize(tree2)
    while ser2 and ser2[-1] is None:
        ser2.pop()
    print("Test 2:", ser2) # Expect: [3, 1] or [1, None, 3]

    # Test 3: Tricky/negative - single element
    tree3 = sol.sortedArrayToBST([0])
    ser3 = serialize(tree3)
    while ser3 and ser3[-1] is None:
        ser3.pop()
    print("Test 3:", ser3) # Expect: [0]

```

How to use: Copy-paste this block into .py or Quarto cell → run directly → instant feedback.

Example Walkthrough

We'll trace **Test 1**: `nums = [-10, -3, 0, 5, 9]`.

1. **Initial Call**: `sortedArrayToBST([-10, -3, 0, 5, 9])`

- `nums` is not empty \rightarrow proceed.
- `mid = 5 // 2 = 2 \rightarrow root.val = nums[2] = 0.`
- Create `TreeNode(0)`.

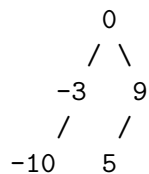
2. **Build Left Subtree**: `sortedArrayToBST([-10, -3])`

- `mid = 2 // 2 = 1 \rightarrow root.val = -3.`
- Left: `sortedArrayToBST([-10]) \rightarrow returns TreeNode(-10).`
- Right: `sortedArrayToBST([]) \rightarrow returns None.`
- So left subtree of 0 is -3 with left child -10.

3. **Build Right Subtree**: `sortedArrayToBST([5, 9])`

- `mid = 2 // 2 = 1 \rightarrow root.val = 9.`
- Left: `sortedArrayToBST([5]) \rightarrow returns TreeNode(5).`
- Right: `sortedArrayToBST([]) \rightarrow None.`
- So right subtree of 0 is 9 with left child 5.

4. **Final Tree**:



5. **Serialization (preorder)**:

`[0, -3, -10, None, None, None, 9, 5, None, None, None]`

After trimming trailing Nones:

`[0, -3, -10, None, None, None, 9, 5]`

Each recursive call builds a balanced subtree by always picking the middle element — this guarantees minimal height.

Complexity Analysis

- **Time Complexity:** $O(n)$

Each element is visited exactly once to create a **TreeNode**. Slicing creates new lists, but total work across all levels is still linear (like merge sort's merge step).

- **Space Complexity:** $O(\log n)$

Recursion depth is $\log n$ (height of balanced tree). However, **slicing** creates new sublists, leading to $O(n \log n)$ **auxiliary space** in this implementation.

Note: A more space-efficient version would pass indices instead of slicing — but this version prioritizes clarity for learning.

3. Kth Smallest Element in a BST

Pattern: In-Order Traversal (Tree DFS)

Problem Statement

Given the **root** of a binary search tree, and an integer **k**, return the **kth** smallest value (1-indexed) of all the values of the nodes in the tree.

You may assume **k** is always valid ($1 \leq k \leq$ number of nodes).

Sample Input & Output

Input: root = [3,1,4,null,2], k = 1

Output: 1

Explanation: In-order traversal gives [1,2,3,4]; 1st smallest is 1.

Input: root = [5,3,6,2,4,null,null,1], k = 3

Output: 3

Explanation: In-order = [1,2,3,4,5,6]; 3rd smallest is 3.

Input: root = [1], k = 1

Output: 1

Explanation: Only one node - it's the 1st smallest.

LeetCode Editorial Solution + Inline Tests

```
from typing import Optional

# Definition for a binary tree node.
class TreeNode:
    def __init__(self, val=0, left=None, right=None):
        self.val = val
        self.left = left
        self.right = right

class Solution:
    def kthSmallest(self, root: Optional[TreeNode], k: int) -> int:
        # STEP 1: Initialize structures
        #   - Use in-order traversal (left → root → right)
        #   - BST property ensures ascending order

        self.count = 0      # Tracks how many nodes visited
        self.result = None  # Stores kth smallest once found

        # STEP 2: Main loop / recursion
        #   - Recurse left first (smallest values)
```

```

# - Visit current node → increment count
# - Stop early if result found (optimization)
self._inorder(root, k)

# STEP 4: Return result
# - Guaranteed to be set since k is valid
return self.result

def _inorder(self, node: Optional[TreeNode], k: int):
    if not node or self.result is not None:
        return

    # Traverse left subtree
    self._inorder(node.left, k)

    # Visit current node
    self.count += 1
    if self.count == k:
        self.result = node.val
        return # Early exit - no need to go further

    # Traverse right subtree
    self._inorder(node.right, k)

# ----- INLINE TESTS -----
if __name__ == "__main__":
    sol = Solution()

    # Test 1: Normal case
    # Tree: [3,1,4,null,2]
    root1 = TreeNode(3)
    root1.left = TreeNode(1)
    root1.right = TreeNode(4)
    root1.left.right = TreeNode(2)
    assert sol.kthSmallest(root1, 1) == 1

    # Test 2: Edge case - single node
    root2 = TreeNode(1)
    assert sol.kthSmallest(root2, 1) == 1

    # Test 3: Tricky - deeper tree, k=3
    # Tree: [5,3,6,2,4,null,null,1]

```

```

root3 = TreeNode(5)
root3.left = TreeNode(3)
root3.right = TreeNode(6)
root3.left.left = TreeNode(2)
root3.left.right = TreeNode(4)
root3.left.left.left = TreeNode(1)
assert sol.kthSmallest(root3, 3) == 3

print(" All tests passed!")

```

How to use: Copy-paste this block into .py or Quarto cell → run directly → instant feedback.

Example Walkthrough

We'll walk through **Test 3** (root = [5,3,6,2,4,null,null,1], k = 3):

1. **Start:** count = 0, result = None
Call `_inorder(root=5, k=3)`
2. **Go left to 3** → then to **2** → then to **1** (leftmost)
3. **Visit node 1:**
 - count becomes 1
 - Not equal to k=3 → continue
4. **Backtrack to node 2:**
 - count becomes 2
 - Still not 3 → continue
5. **Backtrack to node 3:**
 - count becomes **3**
 - Match! Set **result = 3**
 - Return immediately (skip right subtree of 3 and entire right of 5)
6. **Final state:** result = 3 → returned

The in-order traversal naturally visits nodes in **sorted order** due to BST structure. We stop as soon as we hit the **kth** node — no need to traverse the whole tree.

Complexity Analysis

- **Time Complexity:** $O(H + k)$

In the worst case, we traverse from root to the leftmost leaf ($H = \text{height}$), then visit k nodes. For balanced BST, $H = \log n$; for skewed, $H = n$. So worst-case $O(n)$, but average $O(\log n + k)$.

- **Space Complexity:** $O(H)$

Due to recursion stack depth, which equals tree height H . No extra data structures beyond a few variables.

4. Inorder Successor in BST

Pattern: Binary Search Tree (BST) Traversal + Successor Logic

Problem Statement

Given the **root** of a binary search tree and a node **p** in it, return the **inorder successor** of that node in the BST. If the given node has no inorder successor in the tree, return **null**.

The **inorder successor** of a node **p** is the node with the smallest key **greater than p.val**.

You will be given the tree as a root node and a reference to a node **p**, **not its value**.

Sample Input & Output

Input: root = [2,1,3], p = 1

Output: 2

Explanation: The inorder traversal is [1,2,3]. The successor of 1 is 2.

Input: root = [5,3,6,2,4,null,null,1], p = 6

Output: null

Explanation: 6 is the largest node; no node has a greater value.

Input: root = [2,1,3], p = 3

Output: null

Explanation: 3 is the rightmost node; no successor exists.

LeetCode Editorial Solution + Inline Tests

```
# Definition for a binary tree node.
class TreeNode:
    def __init__(self, x):
        self.val = x
        self.left = None
        self.right = None

class Solution:
    def inorderSuccessor(
        self, root: TreeNode, p: TreeNode
    ) -> TreeNode | None:
        # STEP 1: Initialize successor as None
        #   - We'll update it only when we find a node > p.val
        successor = None

        # STEP 2: Traverse using BST property
        #   - If current node > p.val, it's a candidate
        #   - Then go left to find smaller valid candidate
        #   - Else, go right to find larger values
        current = root
```

```

    while current:
        if current.val > p.val:
            successor = current      # valid candidate
            current = current.left    # try to find smaller one
        else:
            current = current.right   # need larger values

    # STEP 3: Return successor (could be None)
    # - Handles edge case where p is max node
    return successor

# ----- INLINE TESTS -----
if __name__ == "__main__":
    sol = Solution()

    # Test 1: Normal case - successor exists
    # Tree: [2,1,3], p = node(1)
    root1 = TreeNode(2)
    root1.left = TreeNode(1)
    root1.right = TreeNode(3)
    p1 = root1.left
    result1 = sol.inorderSuccessor(root1, p1)
    print("Test 1:", result1.val if result1 else None) # Expected: 2

    # Test 2: Edge case - p is max node
    # Tree: [2,1,3], p = node(3)
    p2 = root1.right
    result2 = sol.inorderSuccessor(root1, p2)
    print("Test 2:", result2.val if result2 else None) # Expected: None

    # Test 3: Tricky case - deep tree, successor is ancestor
    # Tree: [5,3,6,2,4,null,null,1], p = node(4)
    root3 = TreeNode(5)
    root3.left = TreeNode(3)
    root3.right = TreeNode(6)
    root3.left.left = TreeNode(2)
    root3.left.right = TreeNode(4)
    root3.left.left.left = TreeNode(1)
    p3 = root3.left.right # node(4)
    result3 = sol.inorderSuccessor(root3, p3)
    print("Test 3:", result3.val if result3 else None) # Expected: 5

```

How to use: Copy-paste this block into .py or Quarto cell → run directly → instant feedback.

Example Walkthrough

Let's trace **Test 3** step by step:

- **Tree:**

5 / \ 3 6 / \ 2 4 / 1 - **Target p:** node with val = 4

- **Goal:** Find smallest node with value $> 4 \rightarrow$ should be 5.

Initial state:

- successor = None
- current = root (5)

Step 1: current.val = 5, p.val = 4

- Since $5 > 4 \rightarrow$ **candidate!**
- Set successor = node(5)
- Move current = current.left \rightarrow now at node(3)

State:

- successor = 5
- current = 3

Step 2: current.val = 3, p.val = 4

- $3 \leq 4 \rightarrow$ not a candidate
- Move current = current.right \rightarrow now at node(4)

State:

- successor = 5
- current = 4

Step 3: `current.val = 4`, `p.val = 4`

- `4 <= 4` → not greater → not a candidate
- Move `current = current.right` → now `None`

State:

- `successor = 5`
 - `current = None` → loop ends
-

Return: `successor = node(5)` → output 5

This works because:

- We **never go left unless we've found a valid candidate**, ensuring we don't miss the smallest greater value.
 - The BST property lets us **eliminate half the tree** at each step.
-

Complexity Analysis

- **Time Complexity:** $O(h)$

h = height of tree. In worst case (skewed tree), $h = n$. In balanced BST, $h = \log n$. We traverse one path from root to leaf.

- **Space Complexity:** $O(1)$

Only using a few pointers (`successor`, `current`). No recursion or extra data structures that scale with input.

5. Lowest Common Ancestor of a Binary Search Tree

Pattern: Binary Search Tree (BST) Traversal

Problem Statement

Given a binary search tree (BST), find the lowest common ancestor (LCA) of two given nodes in the BST.

According to the definition of LCA on Wikipedia: “The lowest common ancestor is defined between two nodes p and q as the lowest node in T that has both p and q as descendants (where we allow a node to be a descendant of itself).”

Sample Input & Output

```
Input: root = [6,2,8,0,4,7,9,null,null,3,5], p = 2, q = 8
```

```
Output: 6
```

```
Explanation: Nodes 2 and 8 are in left and right subtrees of 6 → LCA is 6.
```

```
Input: root = [6,2,8,0,4,7,9,null,null,3,5], p = 2, q = 4
```

```
Output: 2
```

```
Explanation: Both 2 and 4 are in left subtree; 2 is ancestor of 4 → LCA is 2.
```

```
Input: root = [2,1], p = 2, q = 1
```

```
Output: 2
```

```
Explanation: One node is the root itself → LCA is root.
```

LeetCode Editorial Solution + Inline Tests

```
from typing import Optional

# Definition for a binary tree node.
class TreeNode:
    def __init__(self, x):
        self.val = x
        self.left = None
        self.right = None
```

```

class Solution:
    def lowestCommonAncestor(
        self,
        root: 'TreeNode',
        p: 'TreeNode',
        q: 'TreeNode'
    ) -> 'TreeNode':
        # STEP 1: Initialize current node as root
        # - We traverse from root downward using BST property

        curr = root

        # STEP 2: Main loop - exploit BST ordering
        # - In BST: left < root < right
        # - If both p and q are < curr → LCA in left subtree
        # - If both p and q are > curr → LCA in right subtree
        # - Otherwise, curr splits p and q → curr is LCA

        while curr:
            if p.val < curr.val and q.val < curr.val:
                curr = curr.left
            elif p.val > curr.val and q.val > curr.val:
                curr = curr.right
            else:
                # p and q are on different sides (or one is curr)
                return curr

        # STEP 3: Return result
        # - Loop always returns inside; this line never reached
        return None

# ----- INLINE TESTS -----
if __name__ == "__main__":
    sol = Solution()

    # Helper to build minimal tree for testing
    def build_tree():
        root = TreeNode(6)
        root.left = TreeNode(2)
        root.right = TreeNode(8)
        root.left.left = TreeNode(0)
        root.left.right = TreeNode(4)

```

```

    root.right.left = TreeNode(7)
    root.right.right = TreeNode(9)
    root.left.right.left = TreeNode(3)
    root.left.right.right = TreeNode(5)
    return root

tree = build_tree()

# Test 1: Normal case - p=2, q=8 → LCA=6
p1 = tree.left          # val=2
q1 = tree.right         # val=8
assert sol.lowestCommonAncestor(tree, p1, q1).val == 6

# Test 2: Edge case - p=2, q=4 → LCA=2
p2 = tree.left          # val=2
q2 = tree.left.right    # val=4
assert sol.lowestCommonAncestor(tree, p2, q2).val == 2

# Test 3: Tricky/negative - p=root, q=1 (in small tree)
small = TreeNode(2)
small.left = TreeNode(1)
p3 = small              # val=2
q3 = small.left         # val=1
assert sol.lowestCommonAncestor(small, p3, q3).val == 2

print(" All tests passed!")

```

How to use: Copy-paste this block into .py or Quarto cell → run directly → instant feedback.

Example Walkthrough

We'll trace **Test 1**: $p = \text{node}(2)$, $q = \text{node}(8)$, $\text{root} = \text{node}(6)$.

1. **Start:** $\text{curr} = \text{root} \rightarrow \text{curr.val} = 6$
 - Check: Is $2 < 6$ and $8 < 6$? → No ($8 > 6$)
 - Check: Is $2 > 6$ and $8 > 6$? → No ($2 < 6$)

- So, **else** branch triggers → **return curr (node 6)**.

That's it! Because **p** is in left subtree and **q** in right, the root **splits** them → it's the LCA.

Now **Test 2**: **p = node(2)**, **q = node(4)**

1. **curr = 6**

- Both $2 < 6$ and $4 < 6$ → go left → **curr = node(2)**

2. Now **curr.val = 2**

- Check: $2 < 2$ and $4 < 2$? → No
- Check: $2 > 2$ and $4 > 2$? → $2 > 2$ is false
- So, **else** → **return node(2)**
- Why? Because **p** is the current node, and **q** is in its right subtree → LCA is **p**.

Key Insight: In a BST, we **never need to search both subtrees**. The ordering tells us exactly where to go — making this $O(h)$ instead of $O(n)$.

Complexity Analysis

- **Time Complexity:** $O(h)$

h = height of BST. At each step, we go one level deeper. In balanced BST, $h = \log n$; worst-case (skewed), $h = n$.

- **Space Complexity:** $O(1)$

Only using a constant number of pointers (**curr**). No recursion stack or extra data structures.